Analysis of Electrocardiograms via Artificial Neural Networks For a Reliable Assessment of a Possible Myocardial Infarction

Alexey Gavryushin
Guglielmo Pacifico
Juan Fernandez de la Garza
Sunho Kim

Index

1.1 Problem Setting

What is Myocardial Infarction?

The presence of acute myocardial injury, detected by abnormal cardiac biomarkers in the setting of evidence of acute myocardial ischemia.

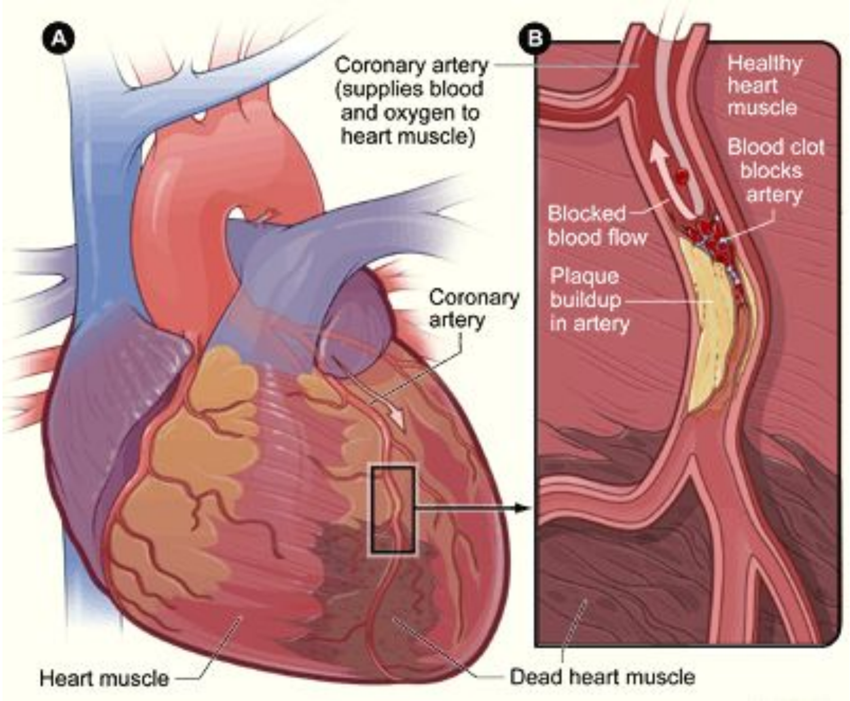

⇩

Happens when one or more areas of the heart muscle don't get enough oxygen. This happens when blood flow to the heart muscle is blocked.

1) https://doi.org/10.1161/CIRCULATIONAHA.119.040631 Circulation. 2019;140:1661–1678
2) https://www.hopkinsmedicine.org/health/conditions-and-diseases/heart-attack

1.1 Problem Setting

Subtypes of Myocardial Infarction

Ischemic?

Deficient supply of blood to a body part that is due to obstruction of the inflow of arterial blood.

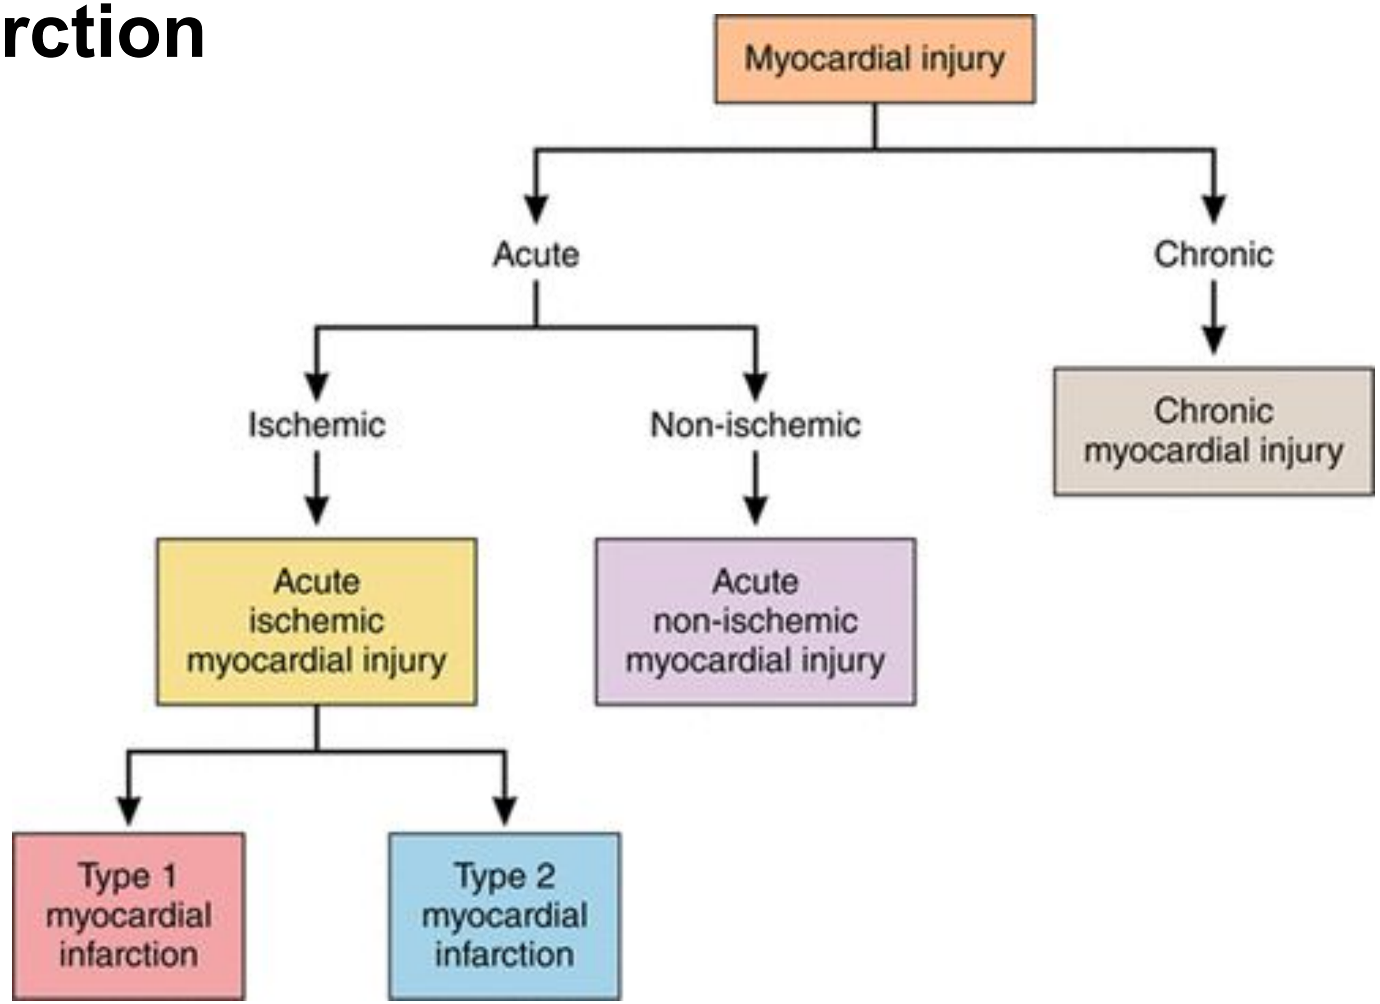

DeFilippis, A. P., Chapman, A. R., Mills, N. L., de Lemos, J. A., Arbab-Zadeh, A., Newby, L. K., & Morrow, D. A. (2019). Assessment and Treatment of Patients With Type 2 Myocardial Infarction and Acute Nonischemic Myocardial Injury. *Circulation*, *140*(20), 1661–1678. https://doi.org/10.1161/CIRCULATIONAHA.119.04063, Figure 1. Myocardial injury taxonomy.

1.1 Problem Setting

How can we detect MI? - ECG data

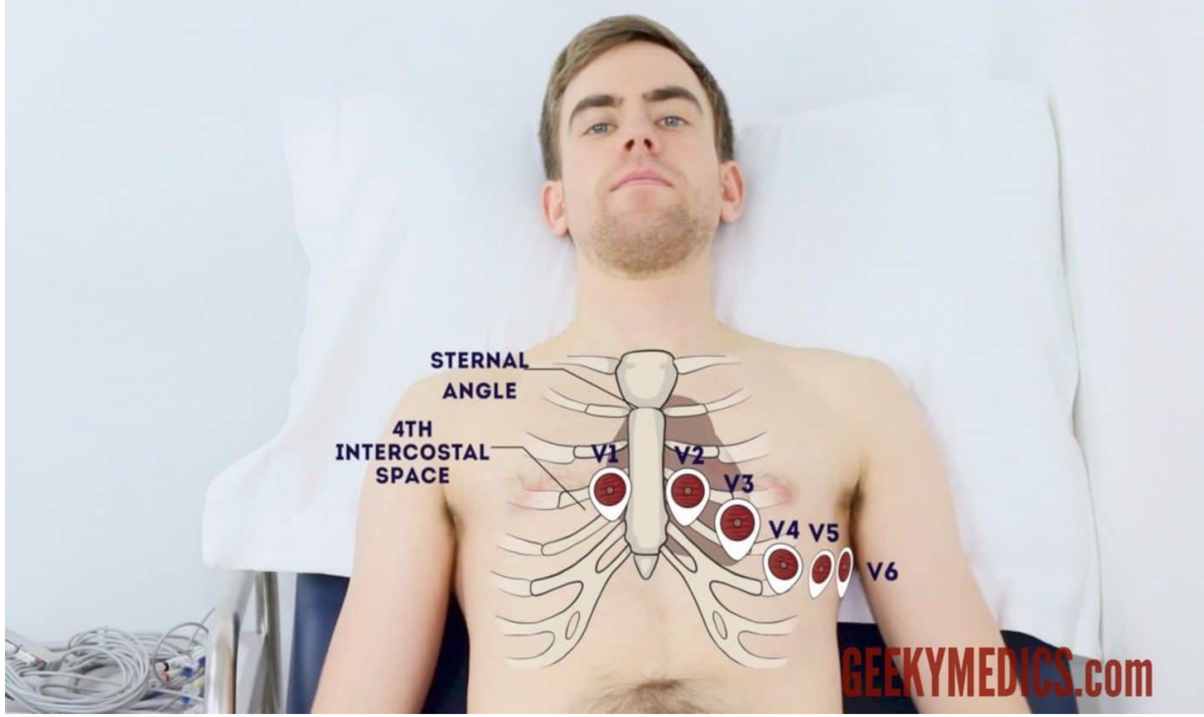

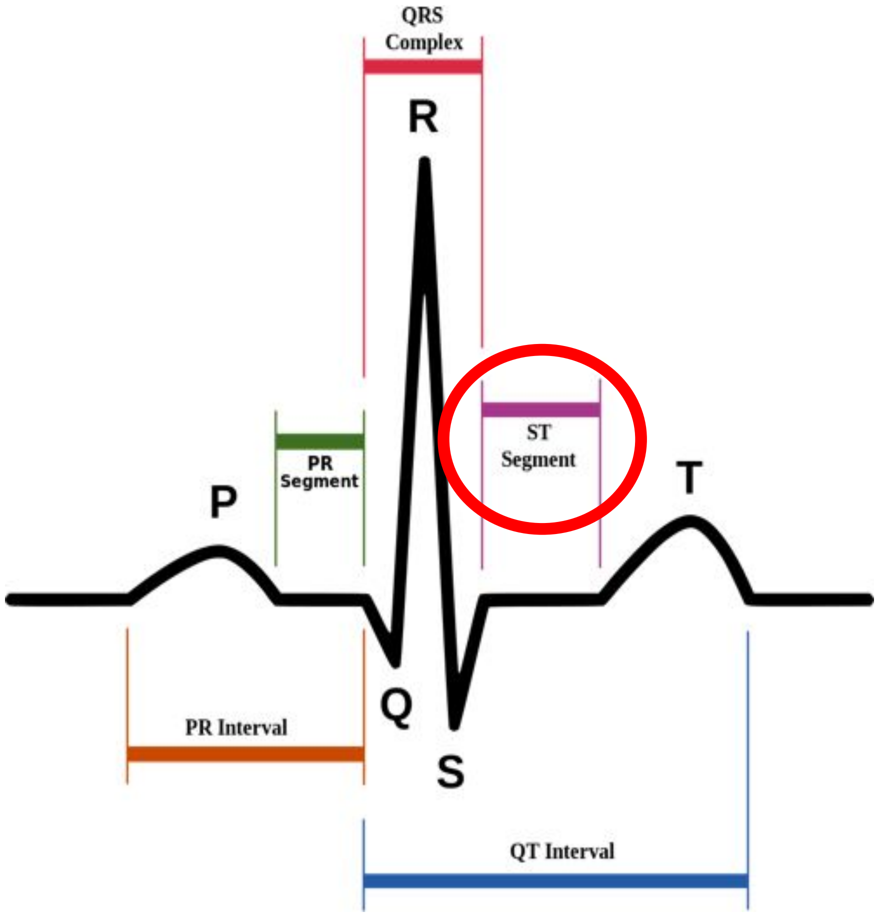

STEMI: ST-elevation myocardial infarction
NSTEMI: heart attack that usually happens when your heart needs oxygen can't be met

https://geekymedics.com/understanding-an-ecg/

1.1 Problem Setting

How can we detect MI?

Figure 1) http://what-when-how.com/paramedic-care/the-monitoring-ecg-clinical-essentials-paramedic-care-part-1
Figure 2) https://www.samsung.com/global/galaxy/galaxy-watch4/
Figure 3) https://www.apple.com/uk/newsroom/2019/03/ecg-app-and-irregular-rhythm-notification-on-apple-watch-available-today-across-europe-and-hong-kong/

1.2 Project Idea

Why do we need a reliable ECG assessment?

"In Tanzania, I was involved in [training and equipping taxi drivers](https://www.hmpgloballearningnetwork.com/site/emsworld/article/12101337/the-struggles-of-ems-in-other-countries) in the place of ambulances and using motorcycle taxi riders as first responders."

Figure 1) https://www.aljazeera.com/gallery/2019/4/1/ugandas-bicycle-ambulances-help-the-pregnant-sick-and-injured
https://www.hmpgloballearningnetwork.com/site/emsworld/article/12101337/the-struggles-of-ems-in-other-countries

2.1 Datasets

- PTB-XL ECG dataset
  - 21837 clinical 12-lead ECGs
  - 18885 patients, between 0 and 95 years old
  - 10 second length
  - 71 total different ECG statements conform to the SCP-ECG s

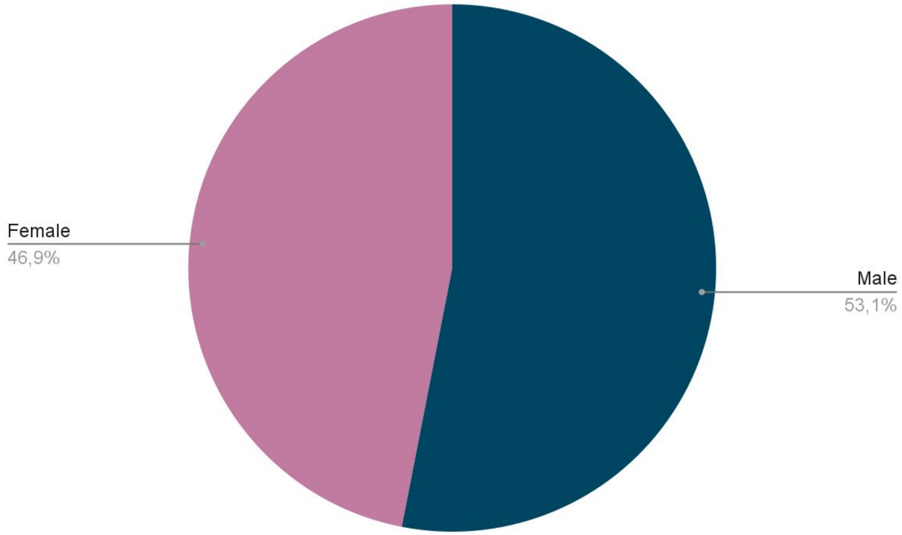

- Generate new “**PTB-V**” (“PTB-Visualized”) dataset
  - From time series to images
  - Python library “PIL” (Pillow) and the classes `Image` and `ImageDraw`

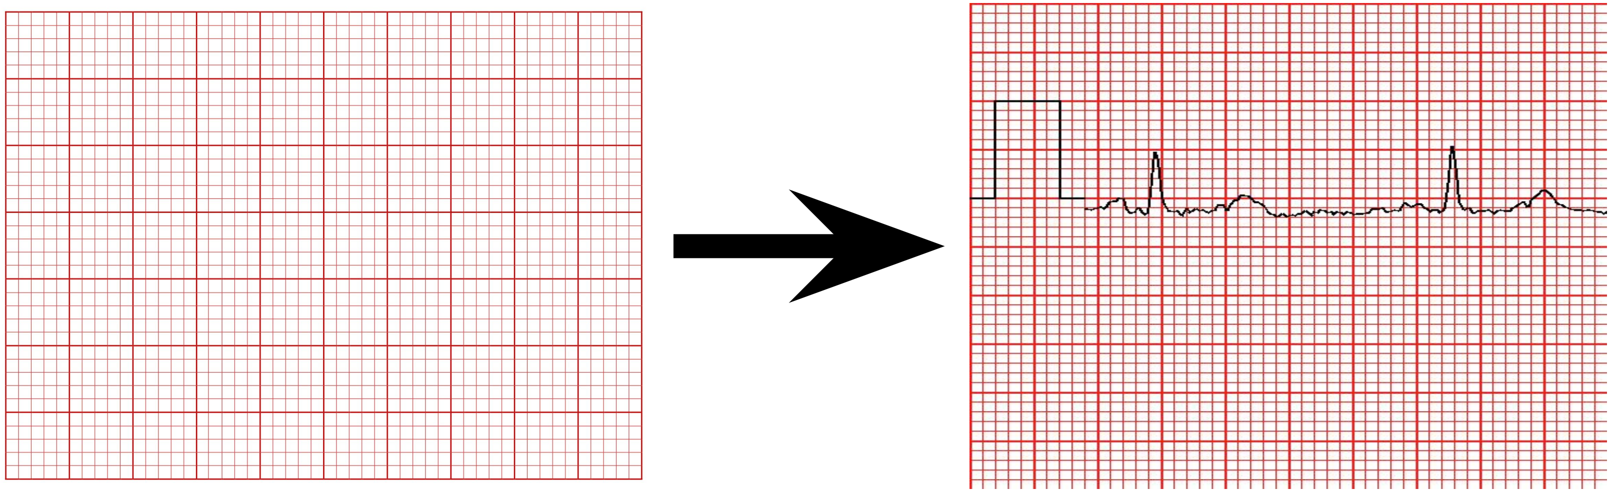

2.2 How to Read an ECG

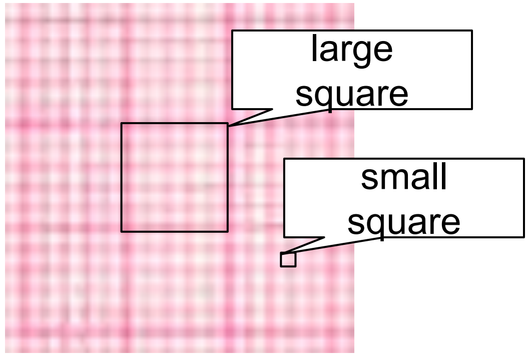

Default paper speed of 25mm/s:

- Each small square is 1 mm^2
- Each small square represents 0.04 s
- Each large square represents 0.2 s
- 5 large squares = 1 s
- 10 squares high = 1 mV

- **Chest leads**
  - V1: septal view of the heart
  - V2: septal view of the heart
  - V3: anterior view of the heart
  - V4: anterior view of the heart
  - V5: lateral view of the heart
  - V6: lateral view of the heart
- **Other leads**
  - Lead I: lateral view
  - Lead II: inferior view
  - Lead III: inferior view
  - aVR: lateral view
  - aVL: lateral view
  - aVF: inferior view

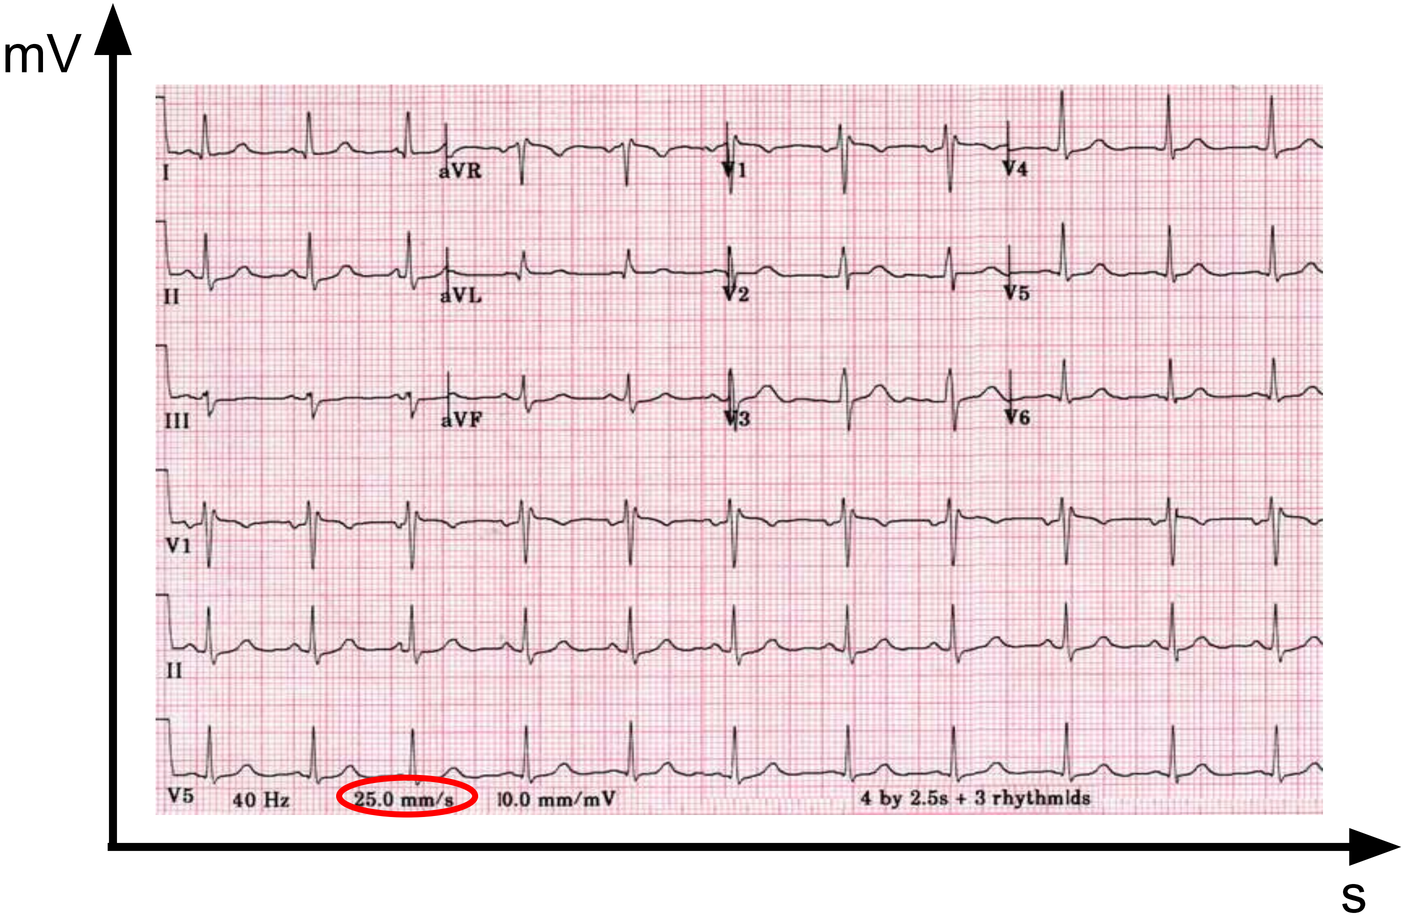

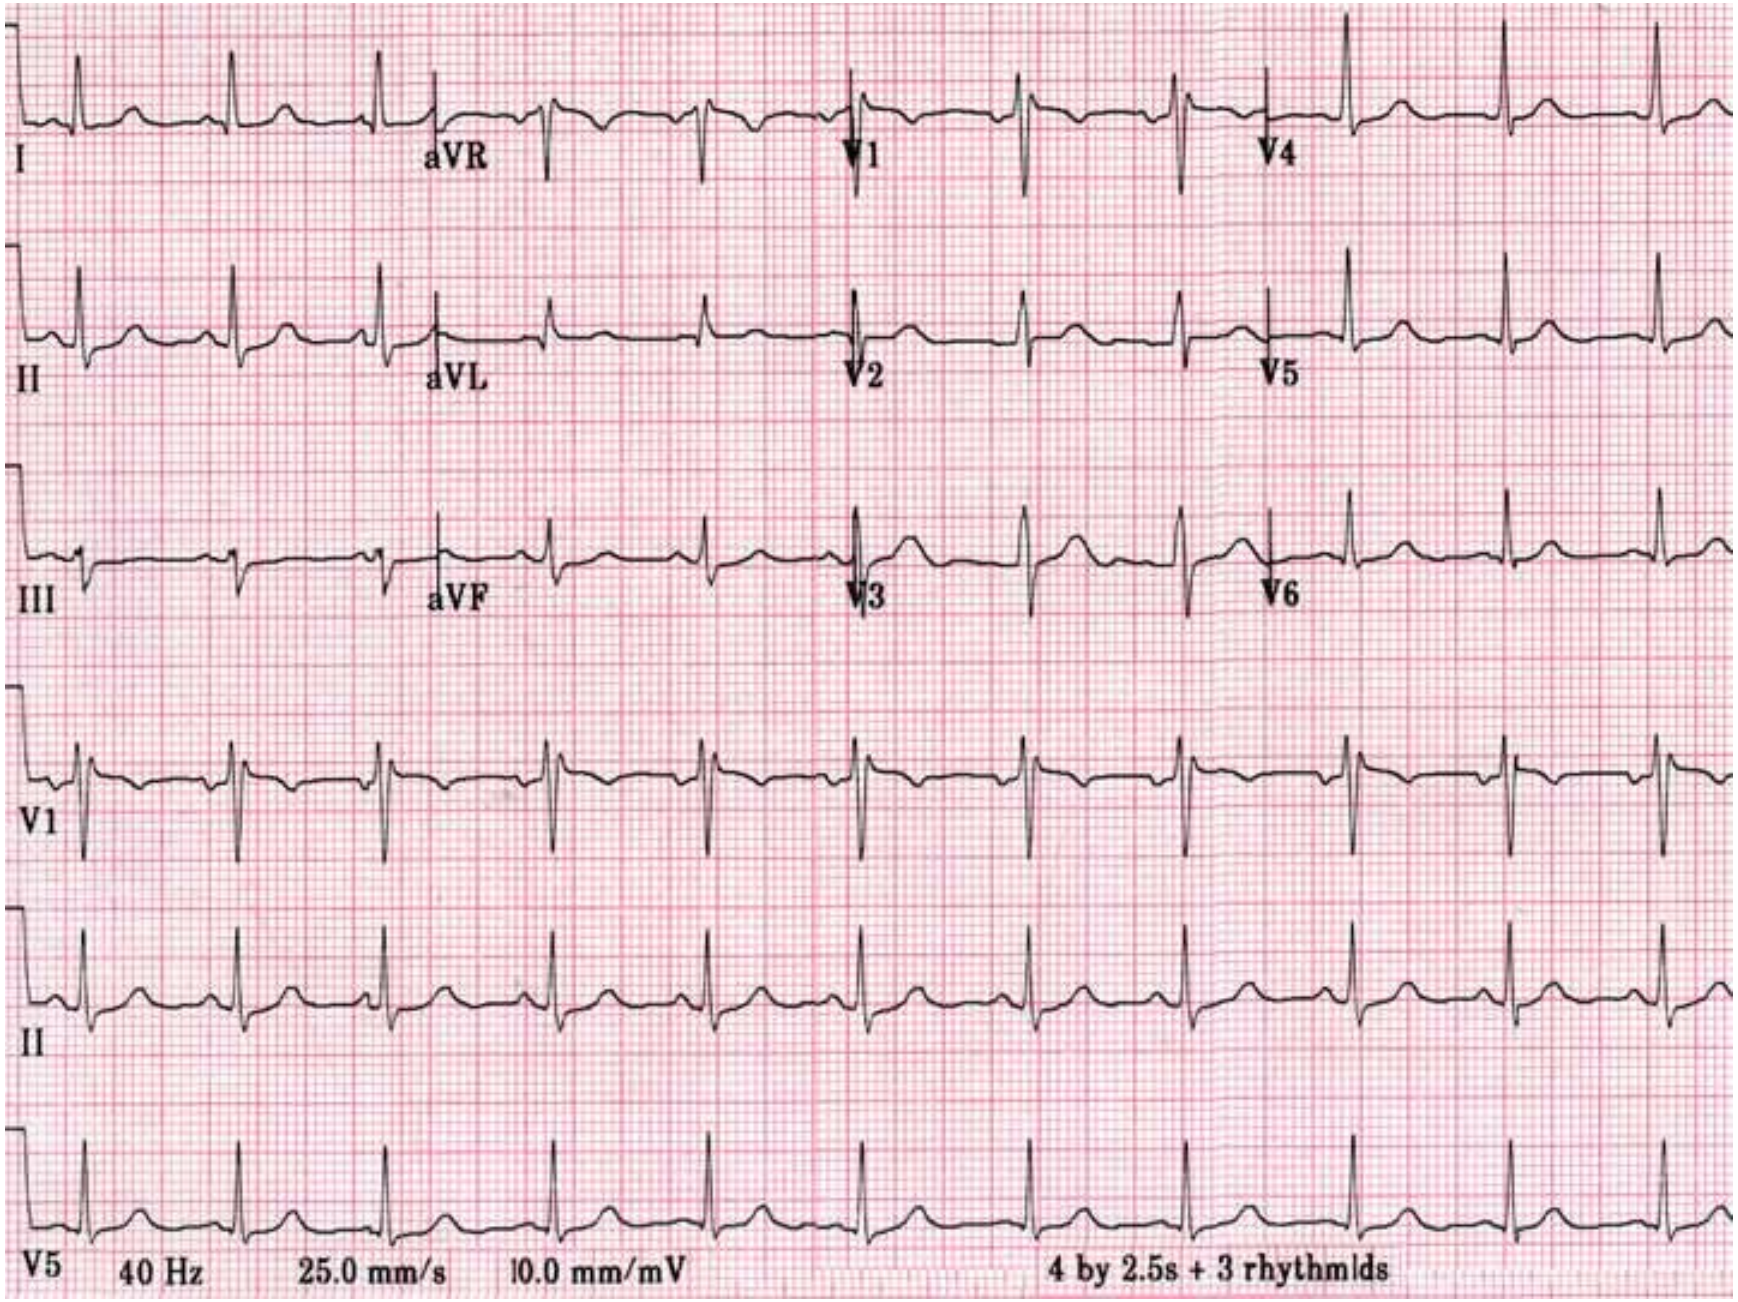
I
aVR
V1
V4
II
aVL
V2
V5
III
aVF
V3
V6
V1
II
V5
40 Hz
25.0 mm/s
10.0 mm/mV
4 by 2.5s + 3 rhythmlds

2.3 Data augmentation

Goal : making the segmentation network more robust to different kinds of inputs

- shearing
- rotation
- blurring
- noise
- shadows and flashes
- different backgrounds

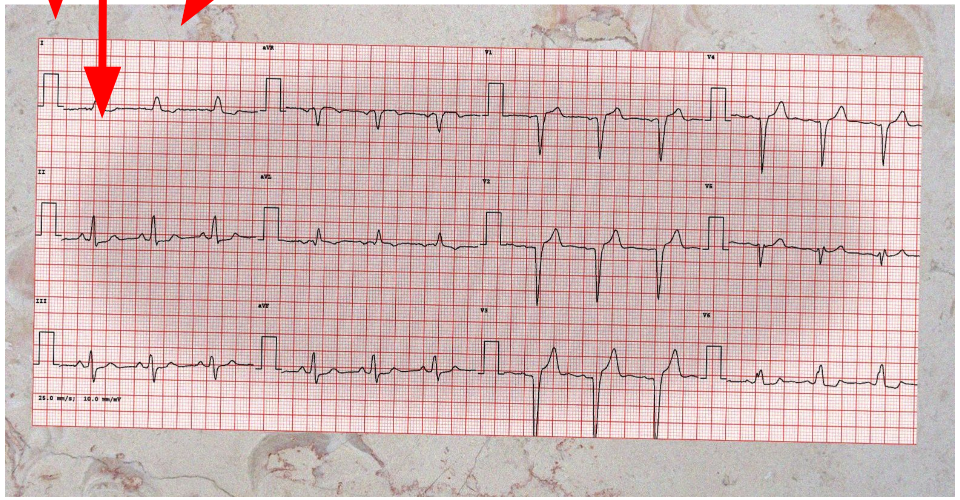

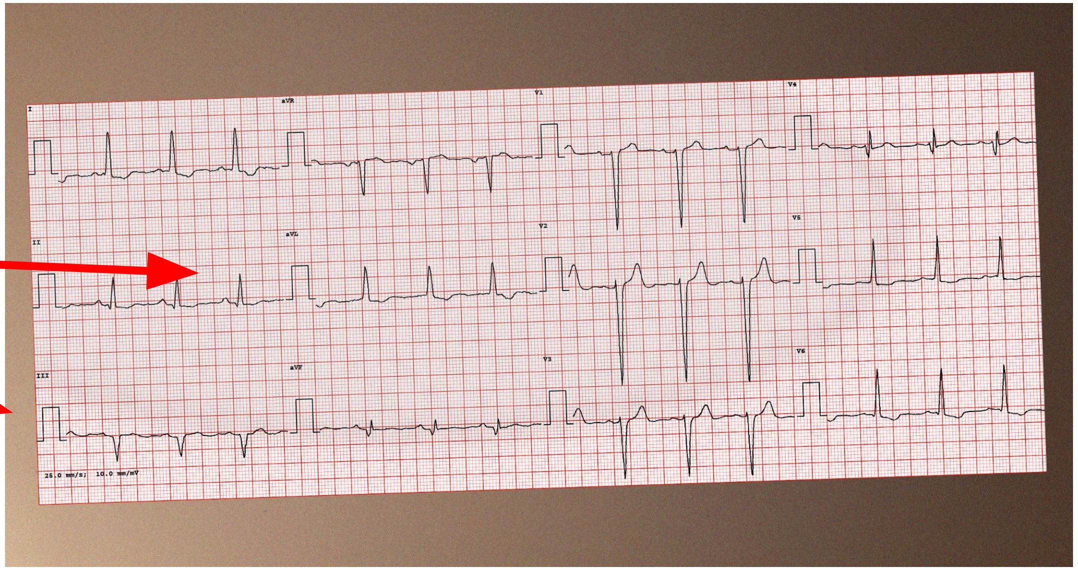

→ as a result we also generate more pictures

2.4 Problems with the data

- **Baseline drift**: Patient moving during the measurement

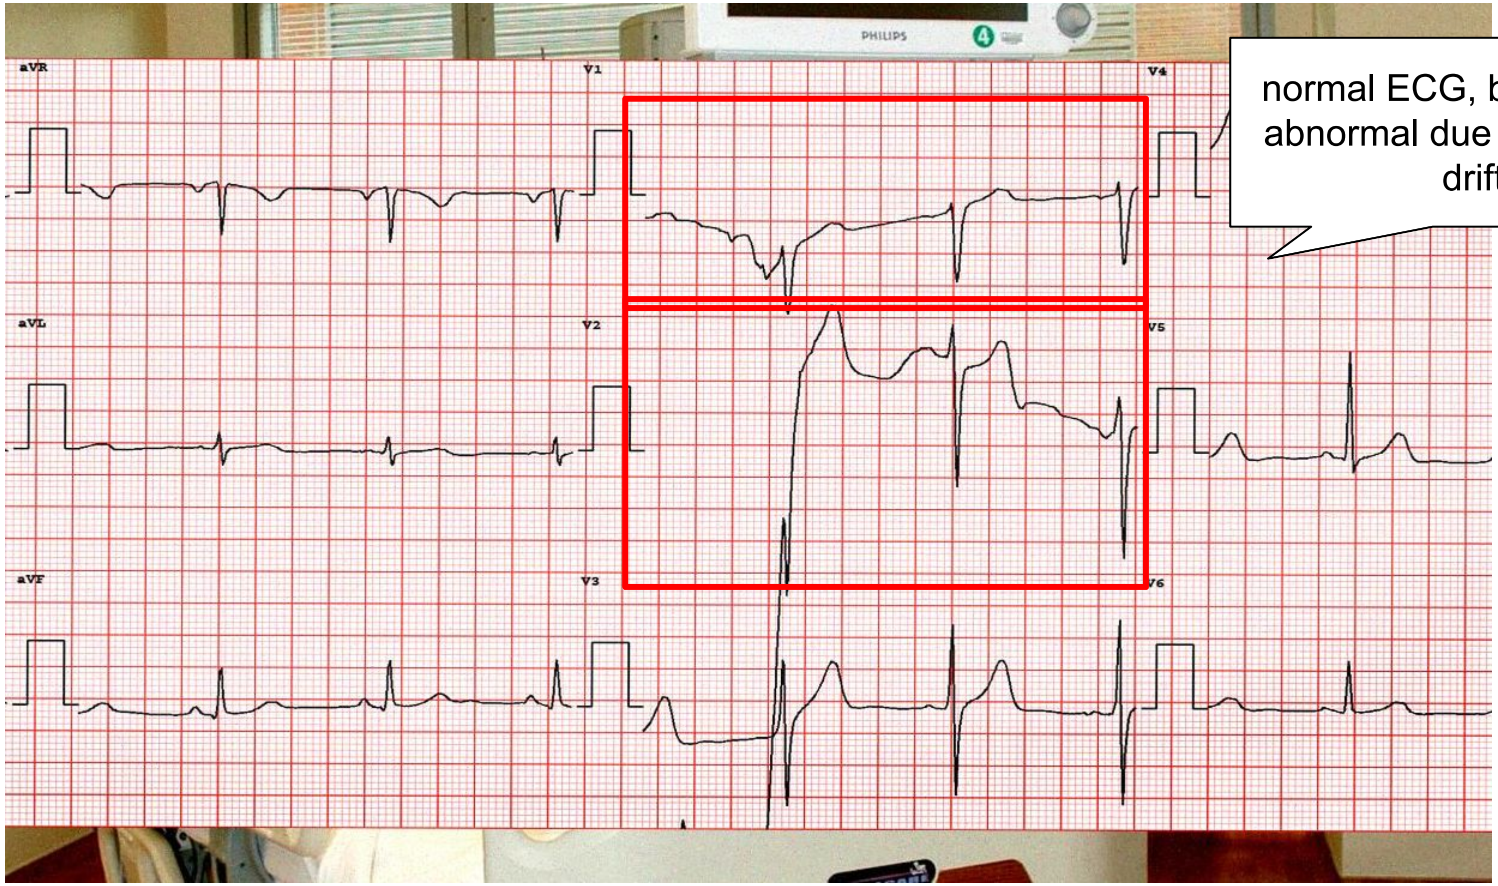

normal ECG, but appears abnormal due to baseline drift

3 ECG Processing Overview

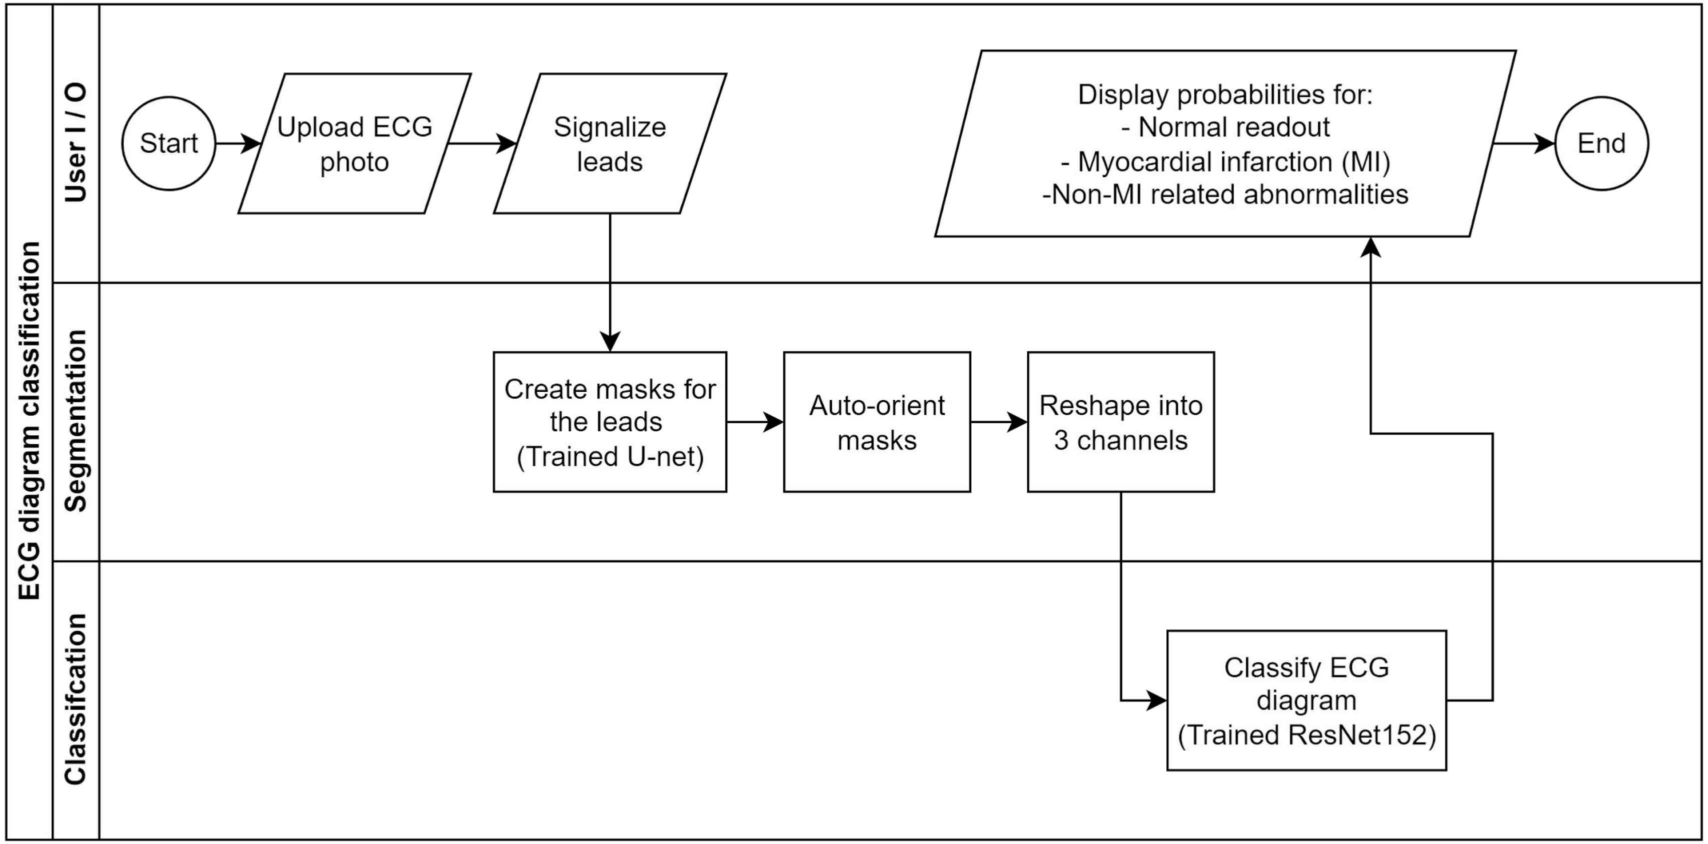

3 ECG Processing Overview

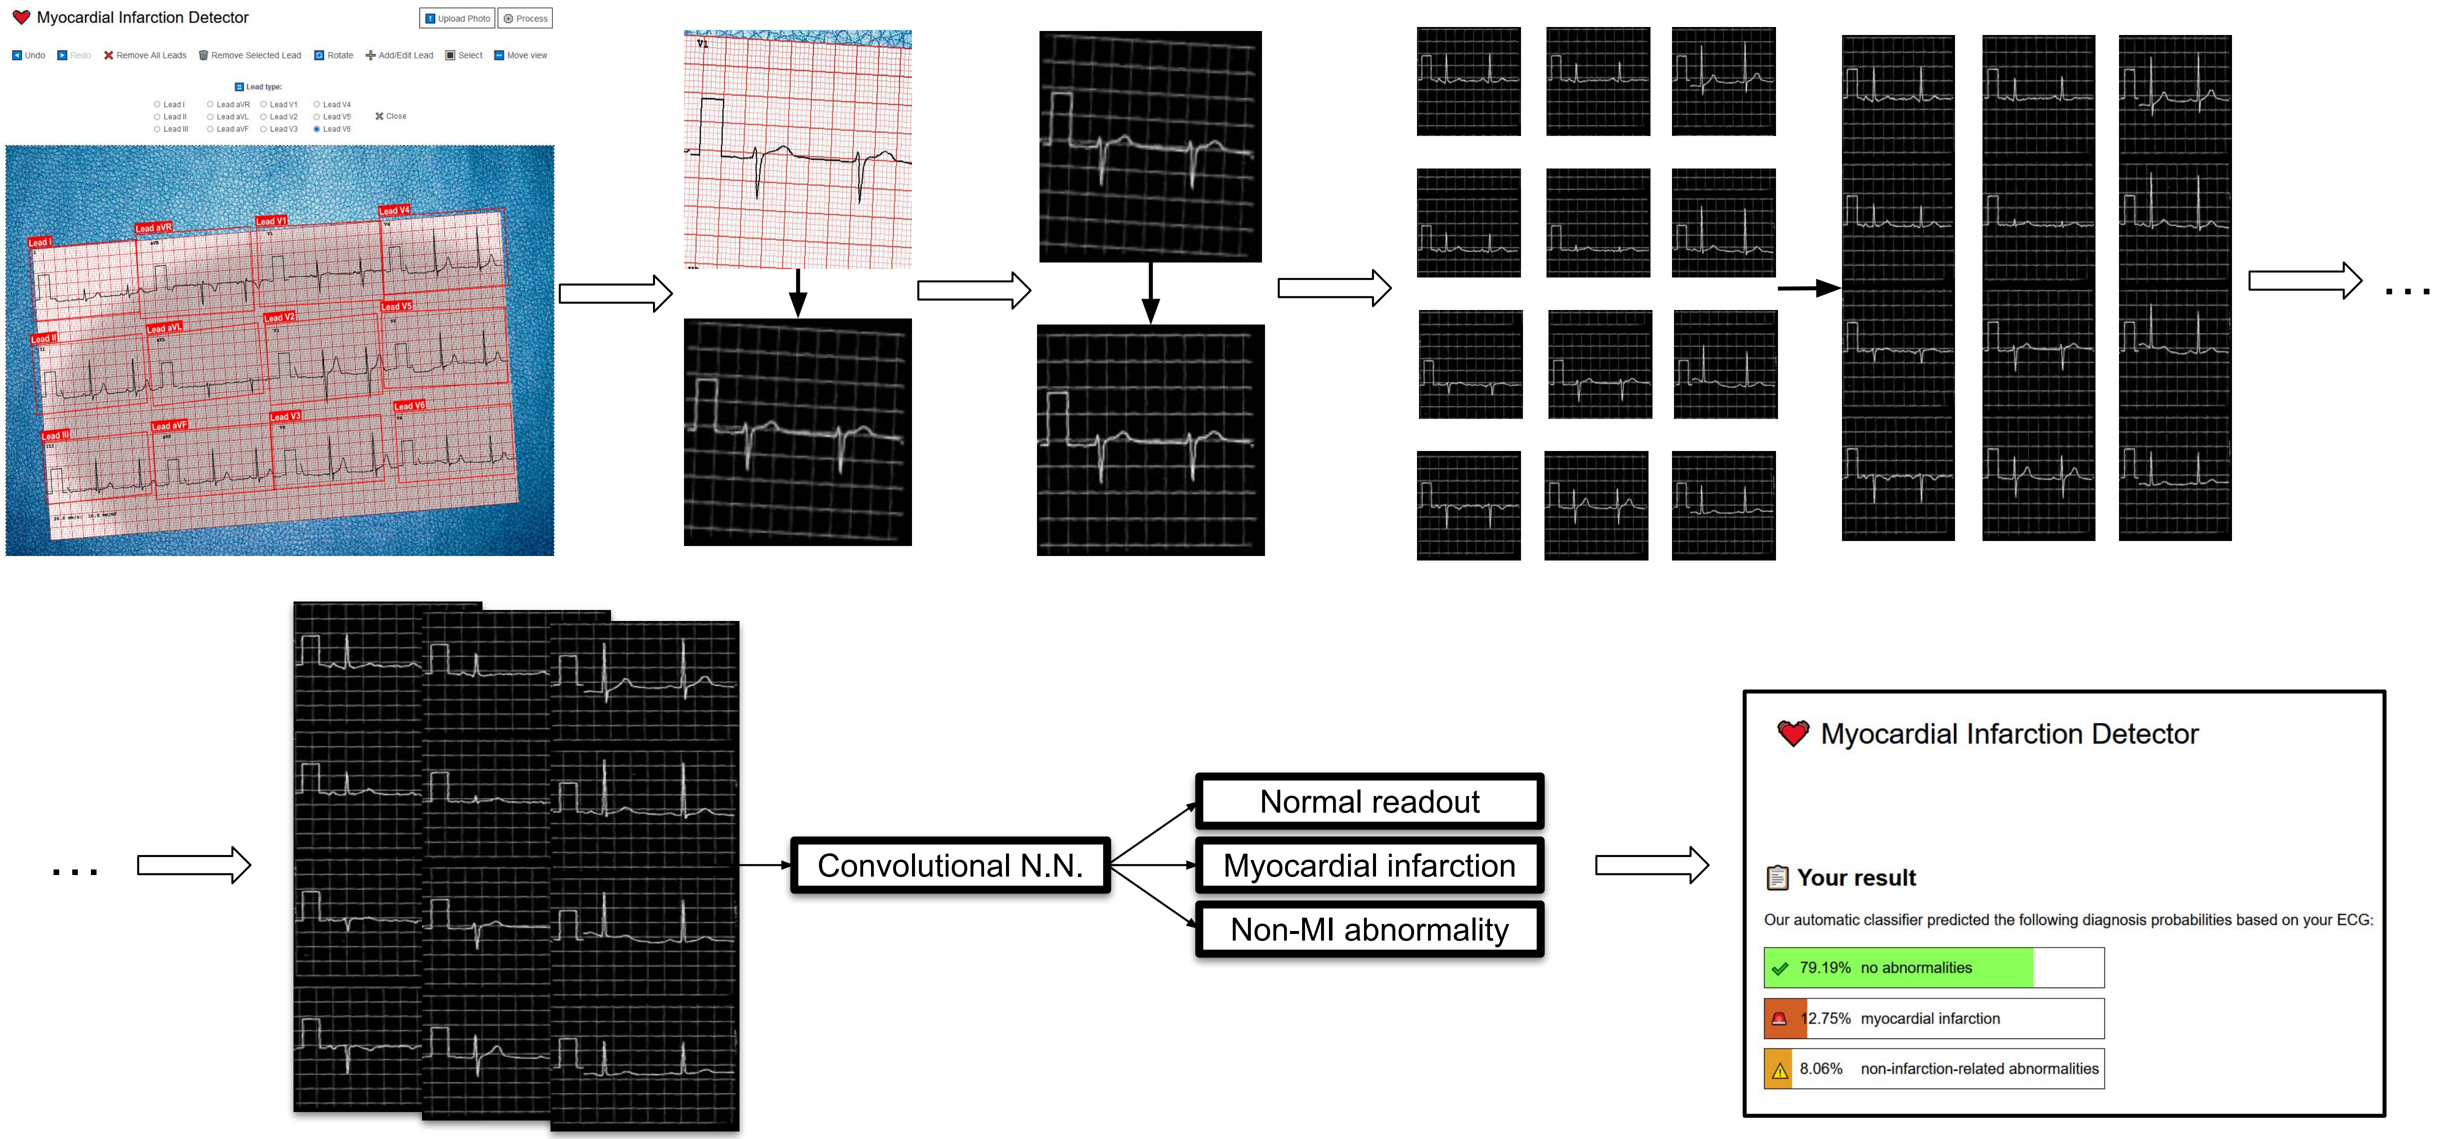

3 ECG Processing Overview

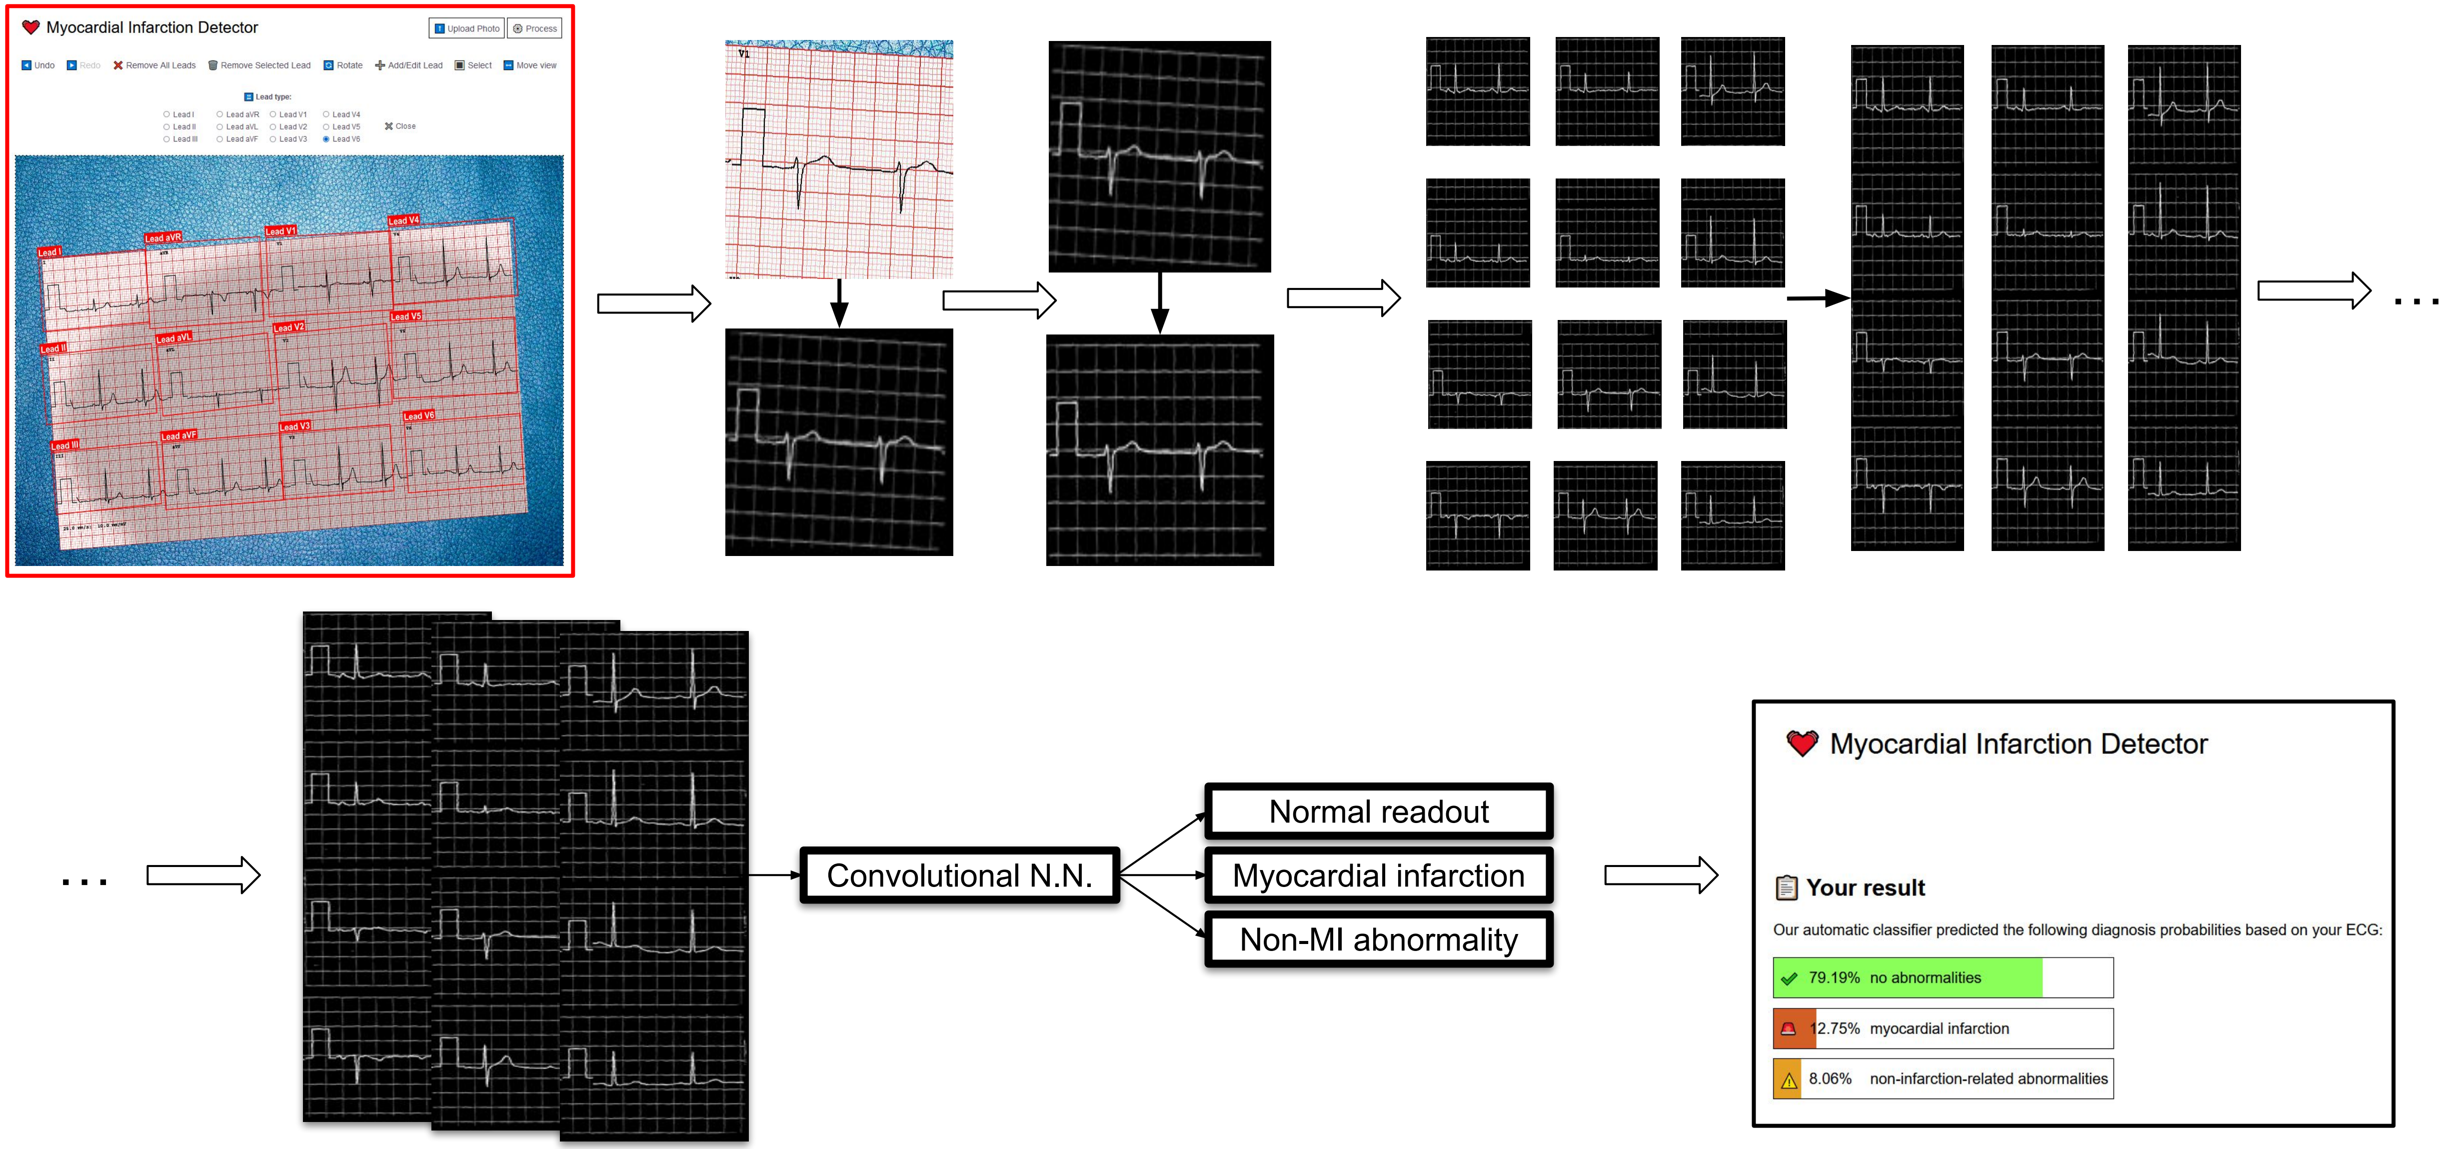

ETH zürich

3 ECG Processing Overview

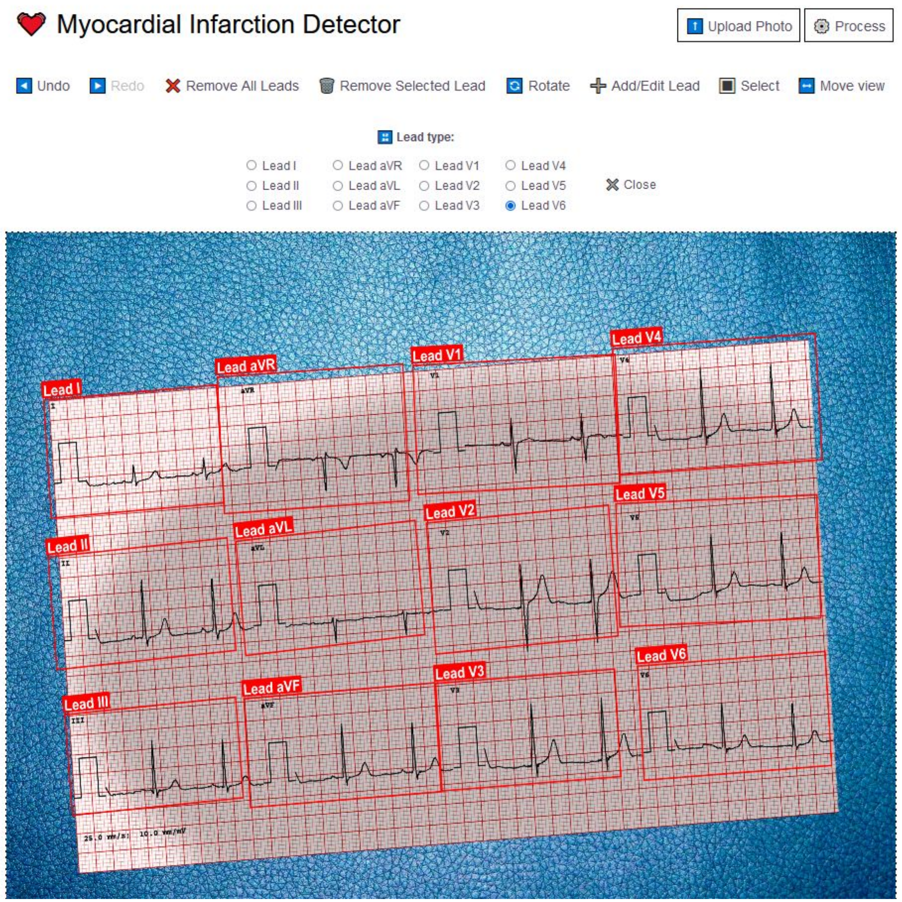

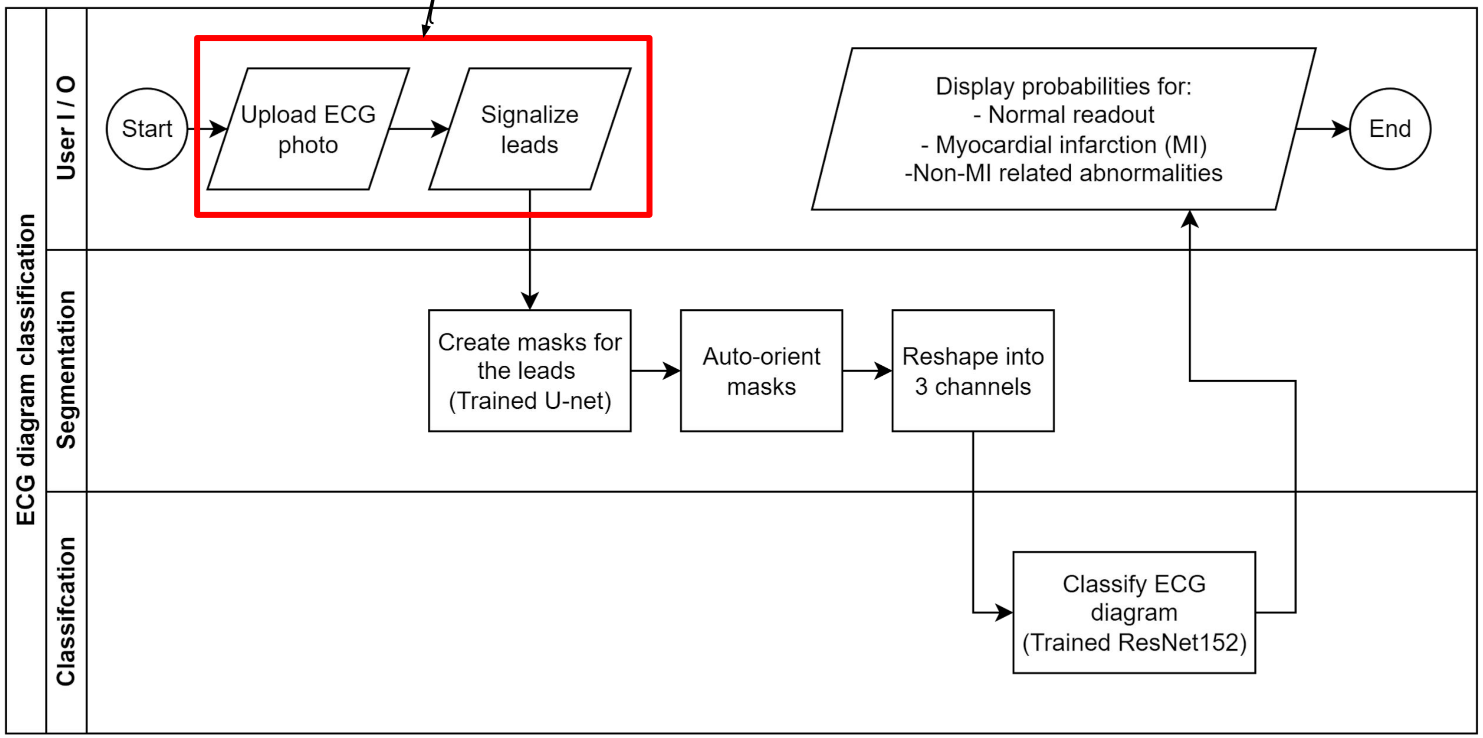

3 ECG Processing Overview

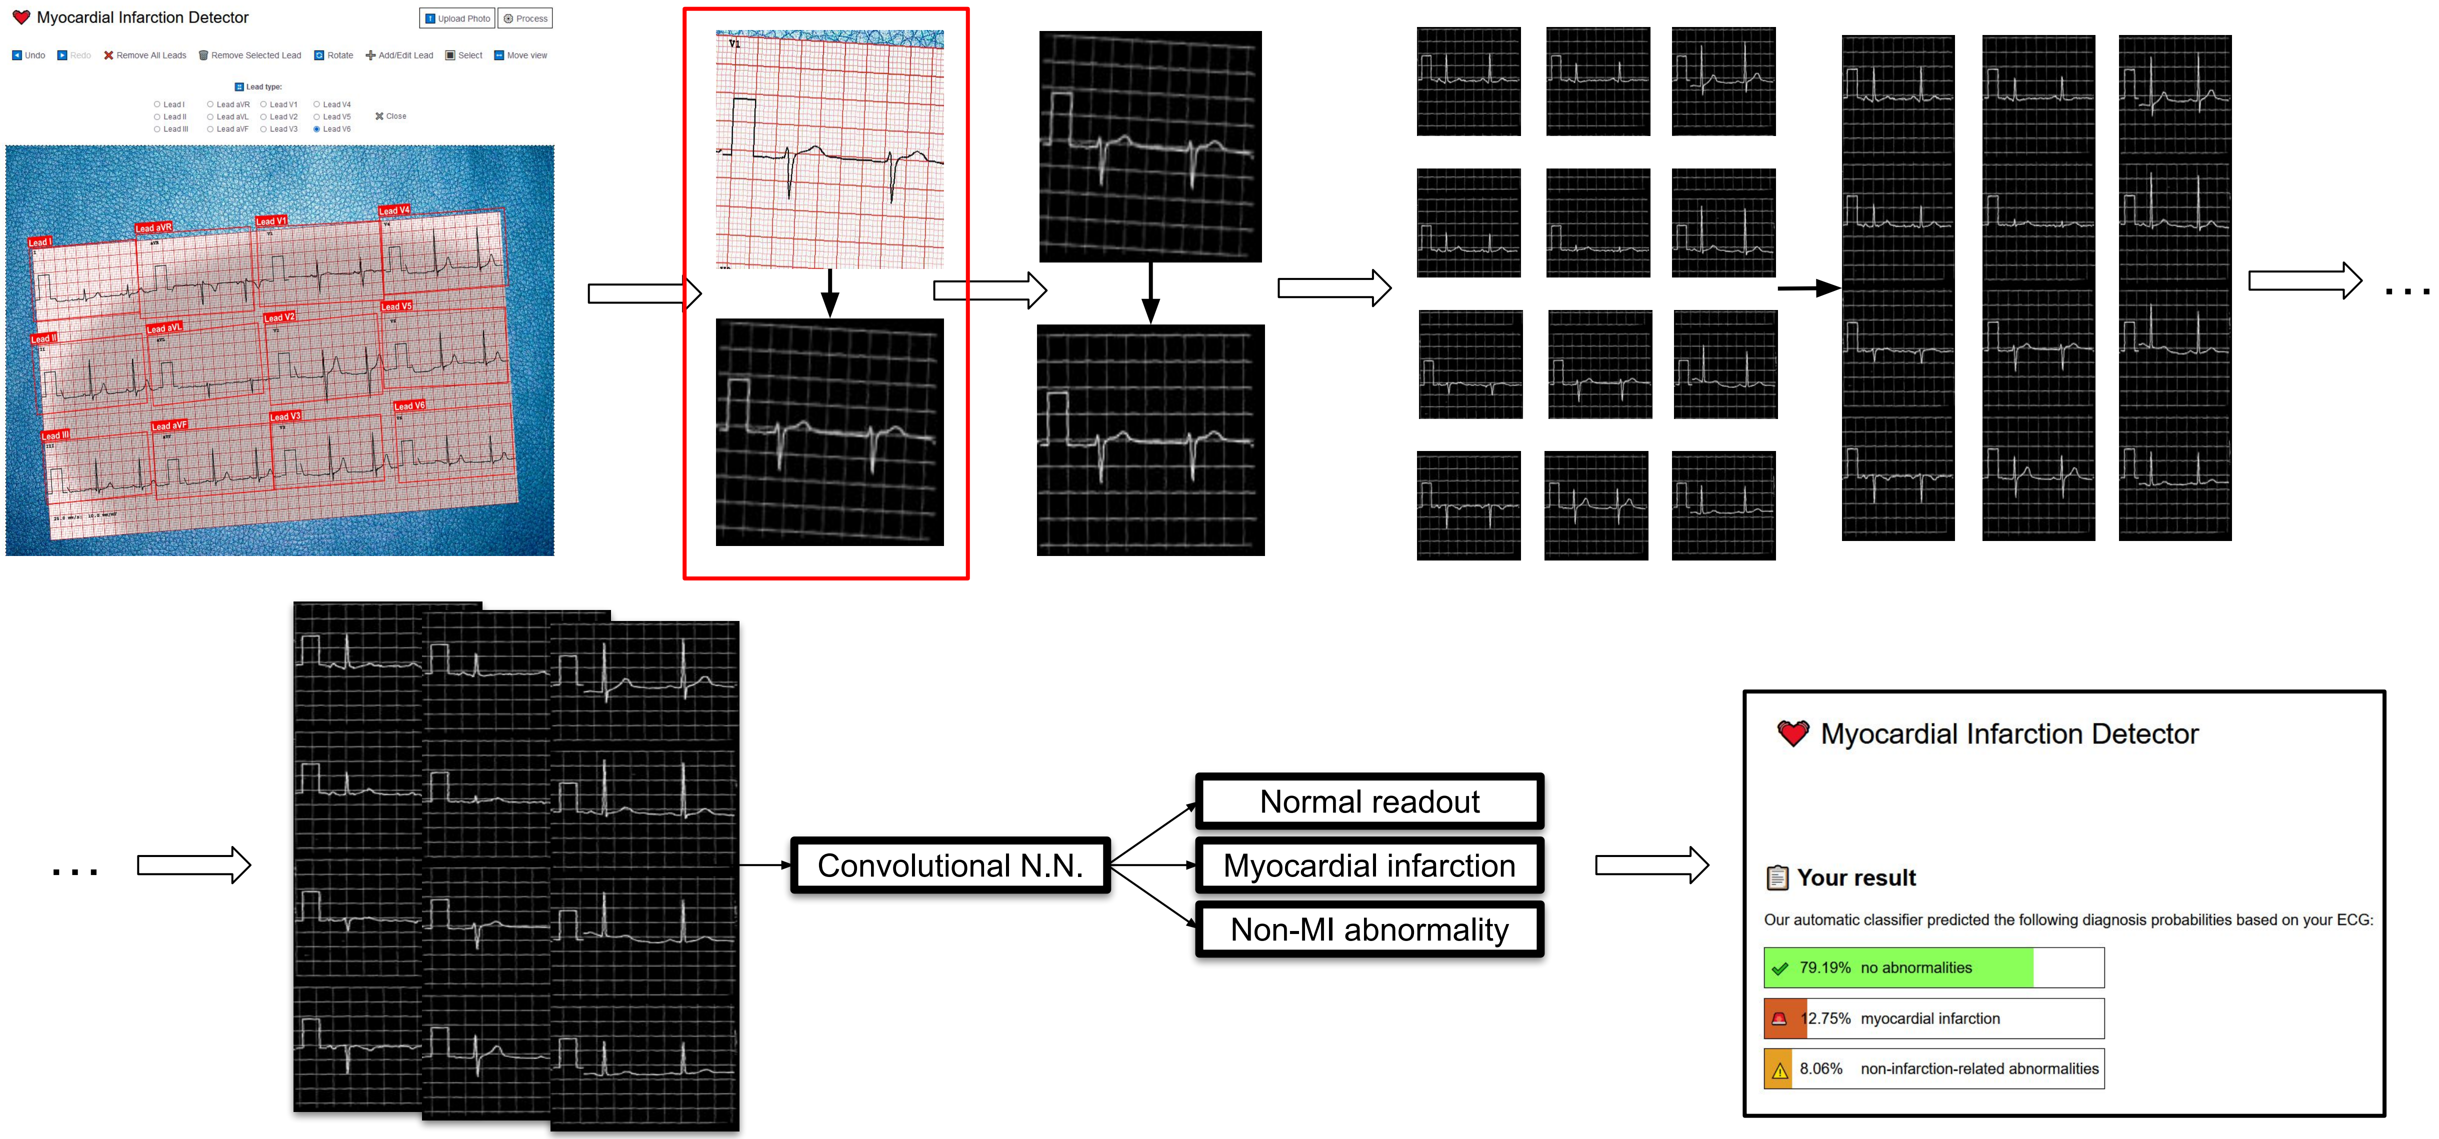

3 ECG Processing Overview

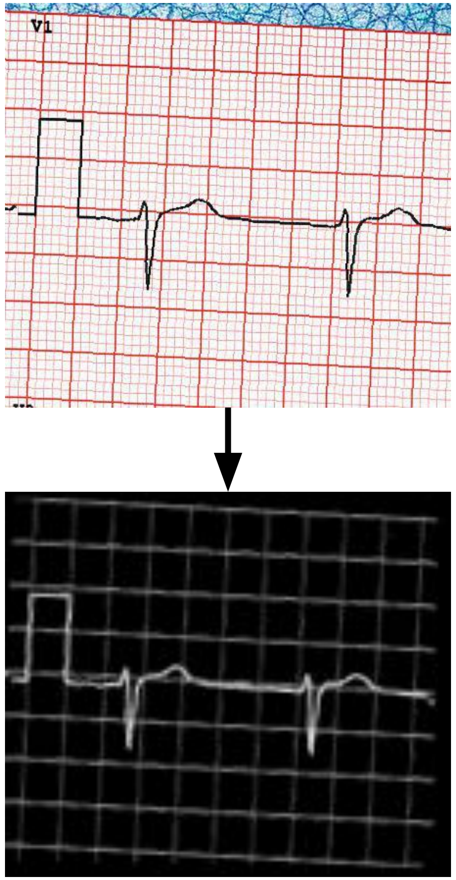

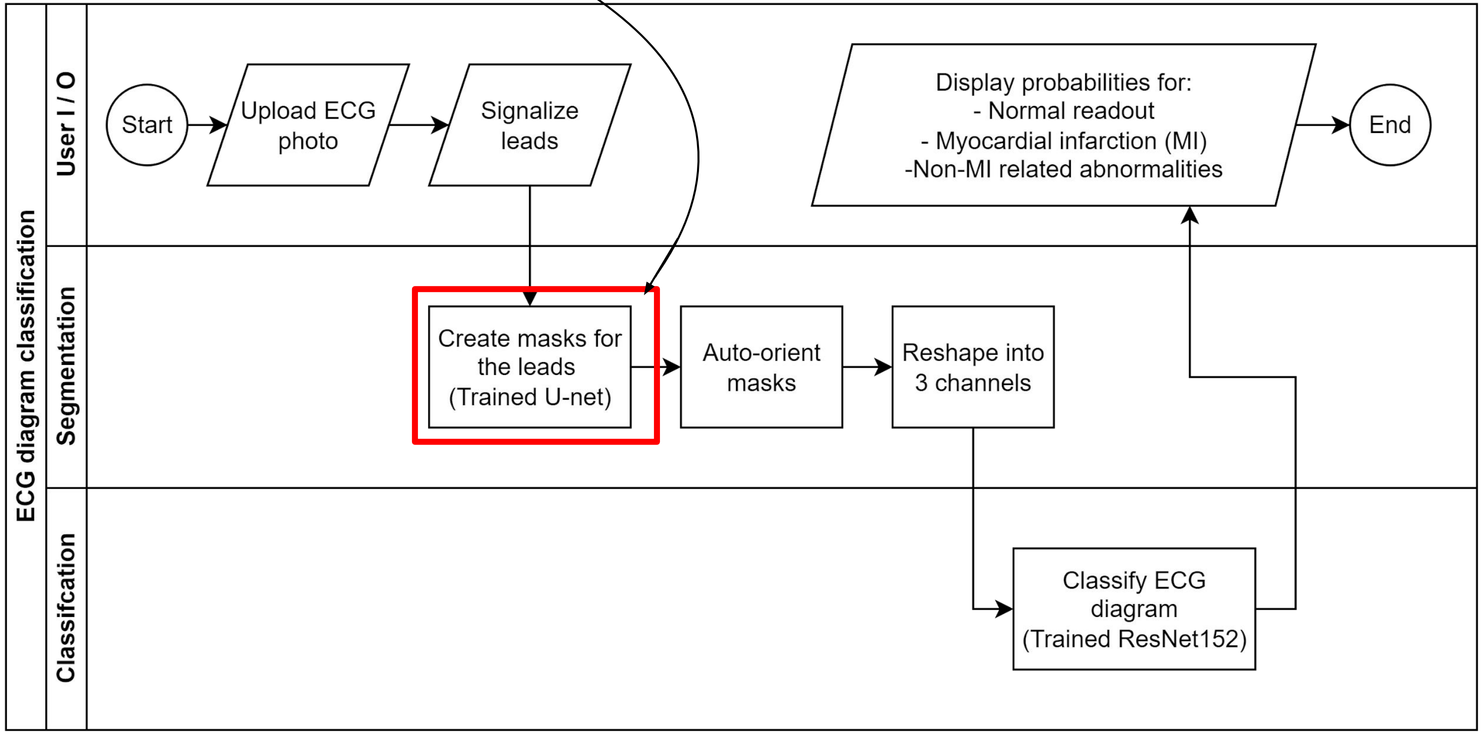

Ronneberger, O., Fischer, P., & Brox, T.. (2015). U-Net: Convolutional Networks for Biomedical Image Segmentation. : https://arxiv.org/abs/1505.04597

3.1 ECG Segmentation

The U-Net was trained with pairs of masks and generated ECG diagrams with data augmentation applied to them, notice that the mask corresponds to the training label.

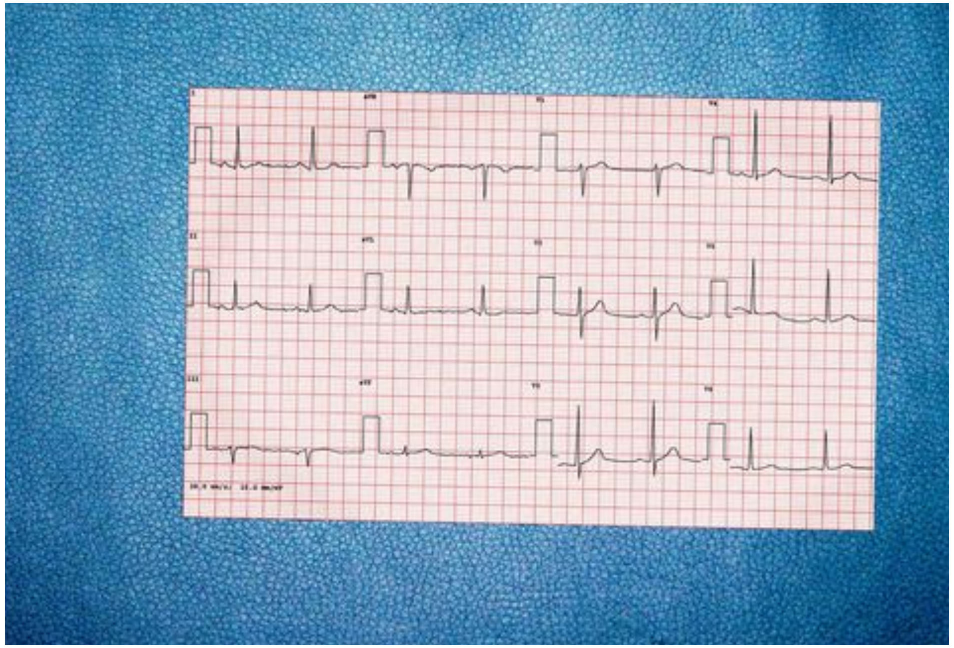

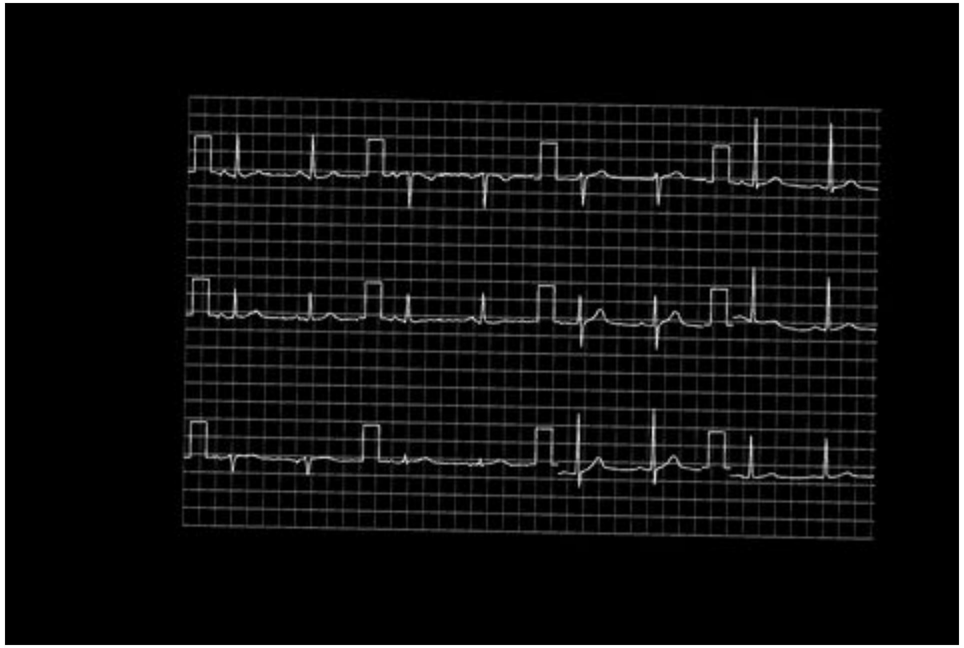

Grid-lines are also captured in the training to be used for auto-orienting of diagrams during processing.

3.1 ECG Segmentation

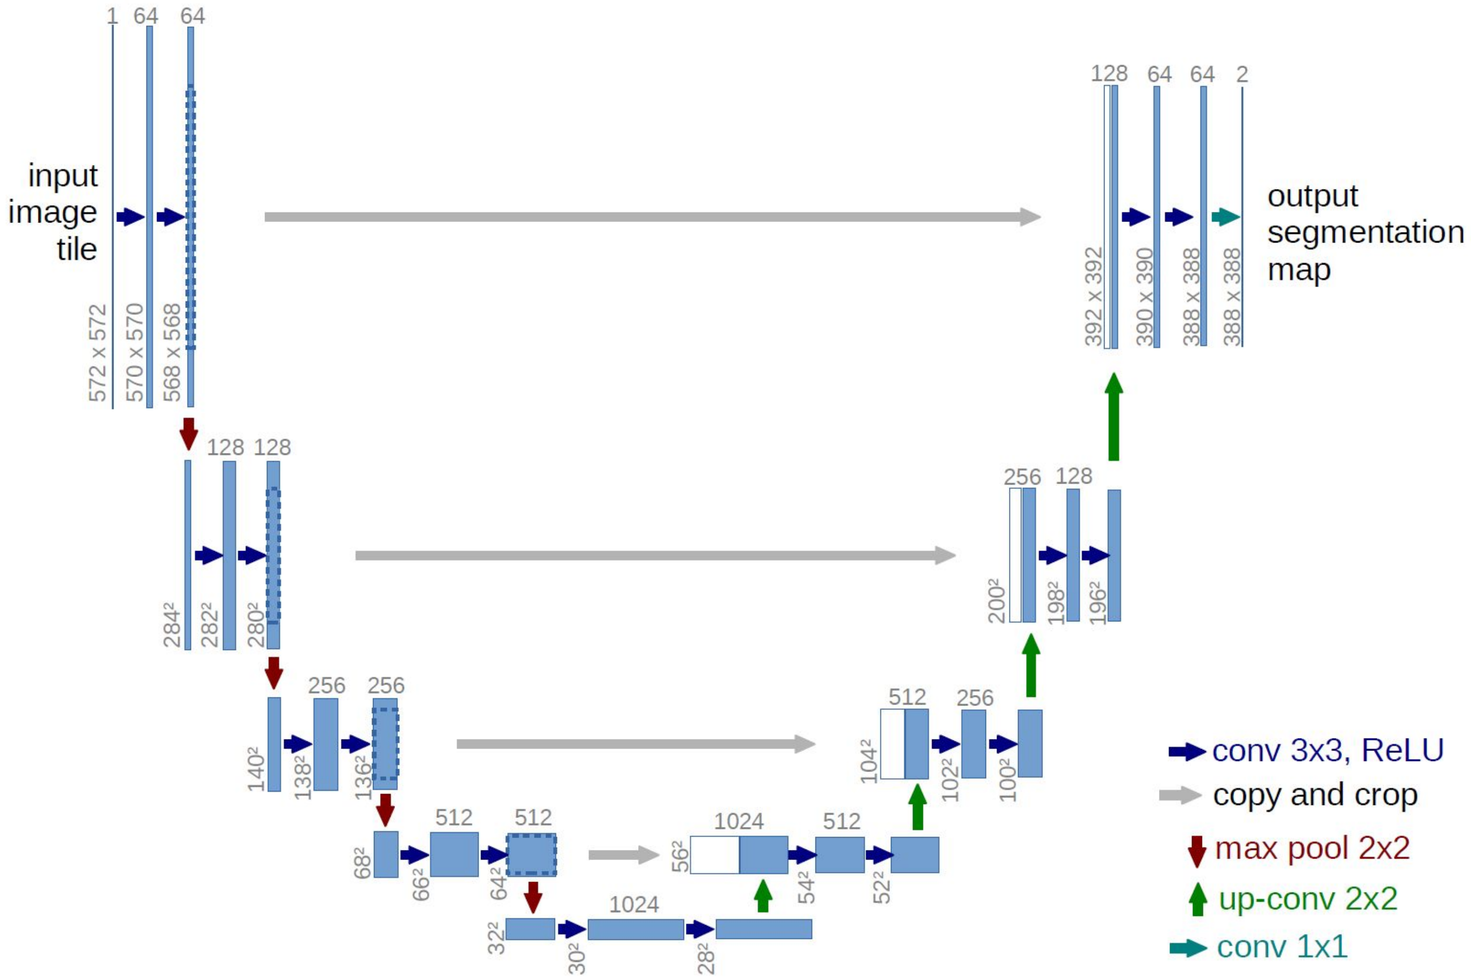

Figure taken from: Ronneberger, O., Fischer, P., & Brox, T.. (2015). U-Net: Convolutional Networks for Biomedical Image Segmentation. : https://arxiv.org/abs/1505.04597

3.1 ECG Segmentation

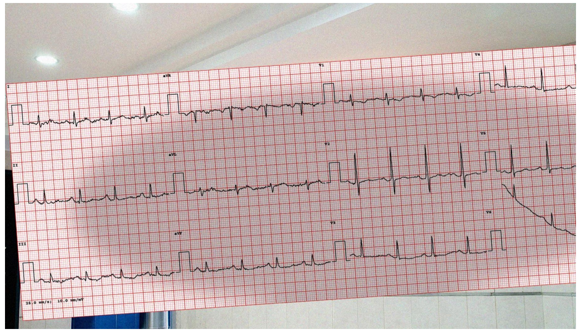

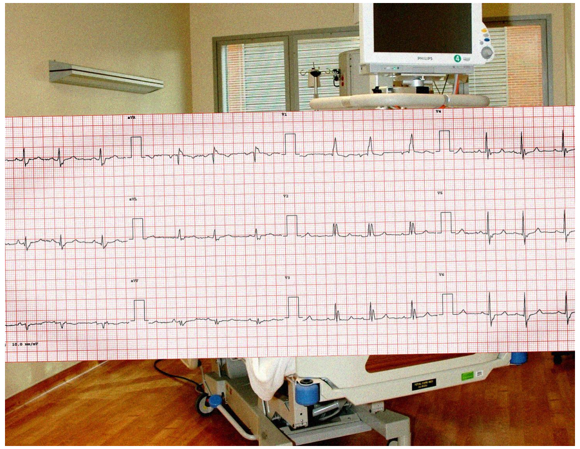
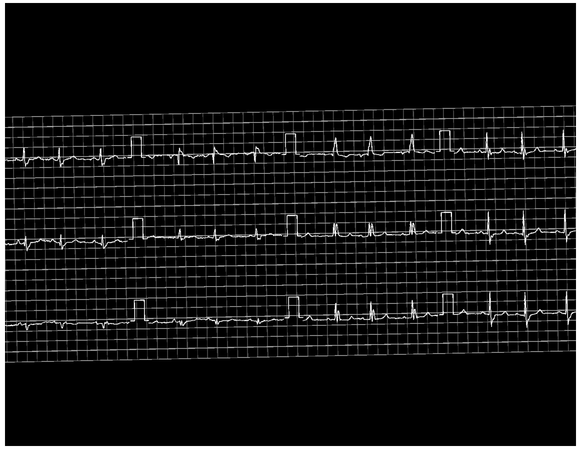
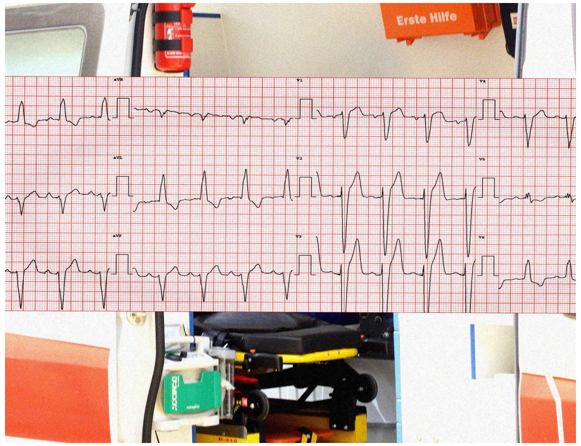


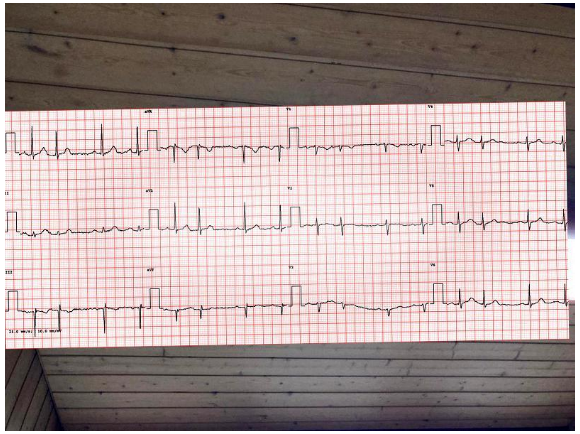
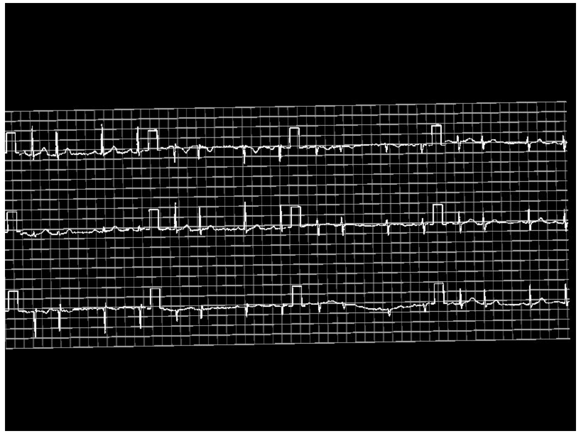

3 ECG Processing Overview

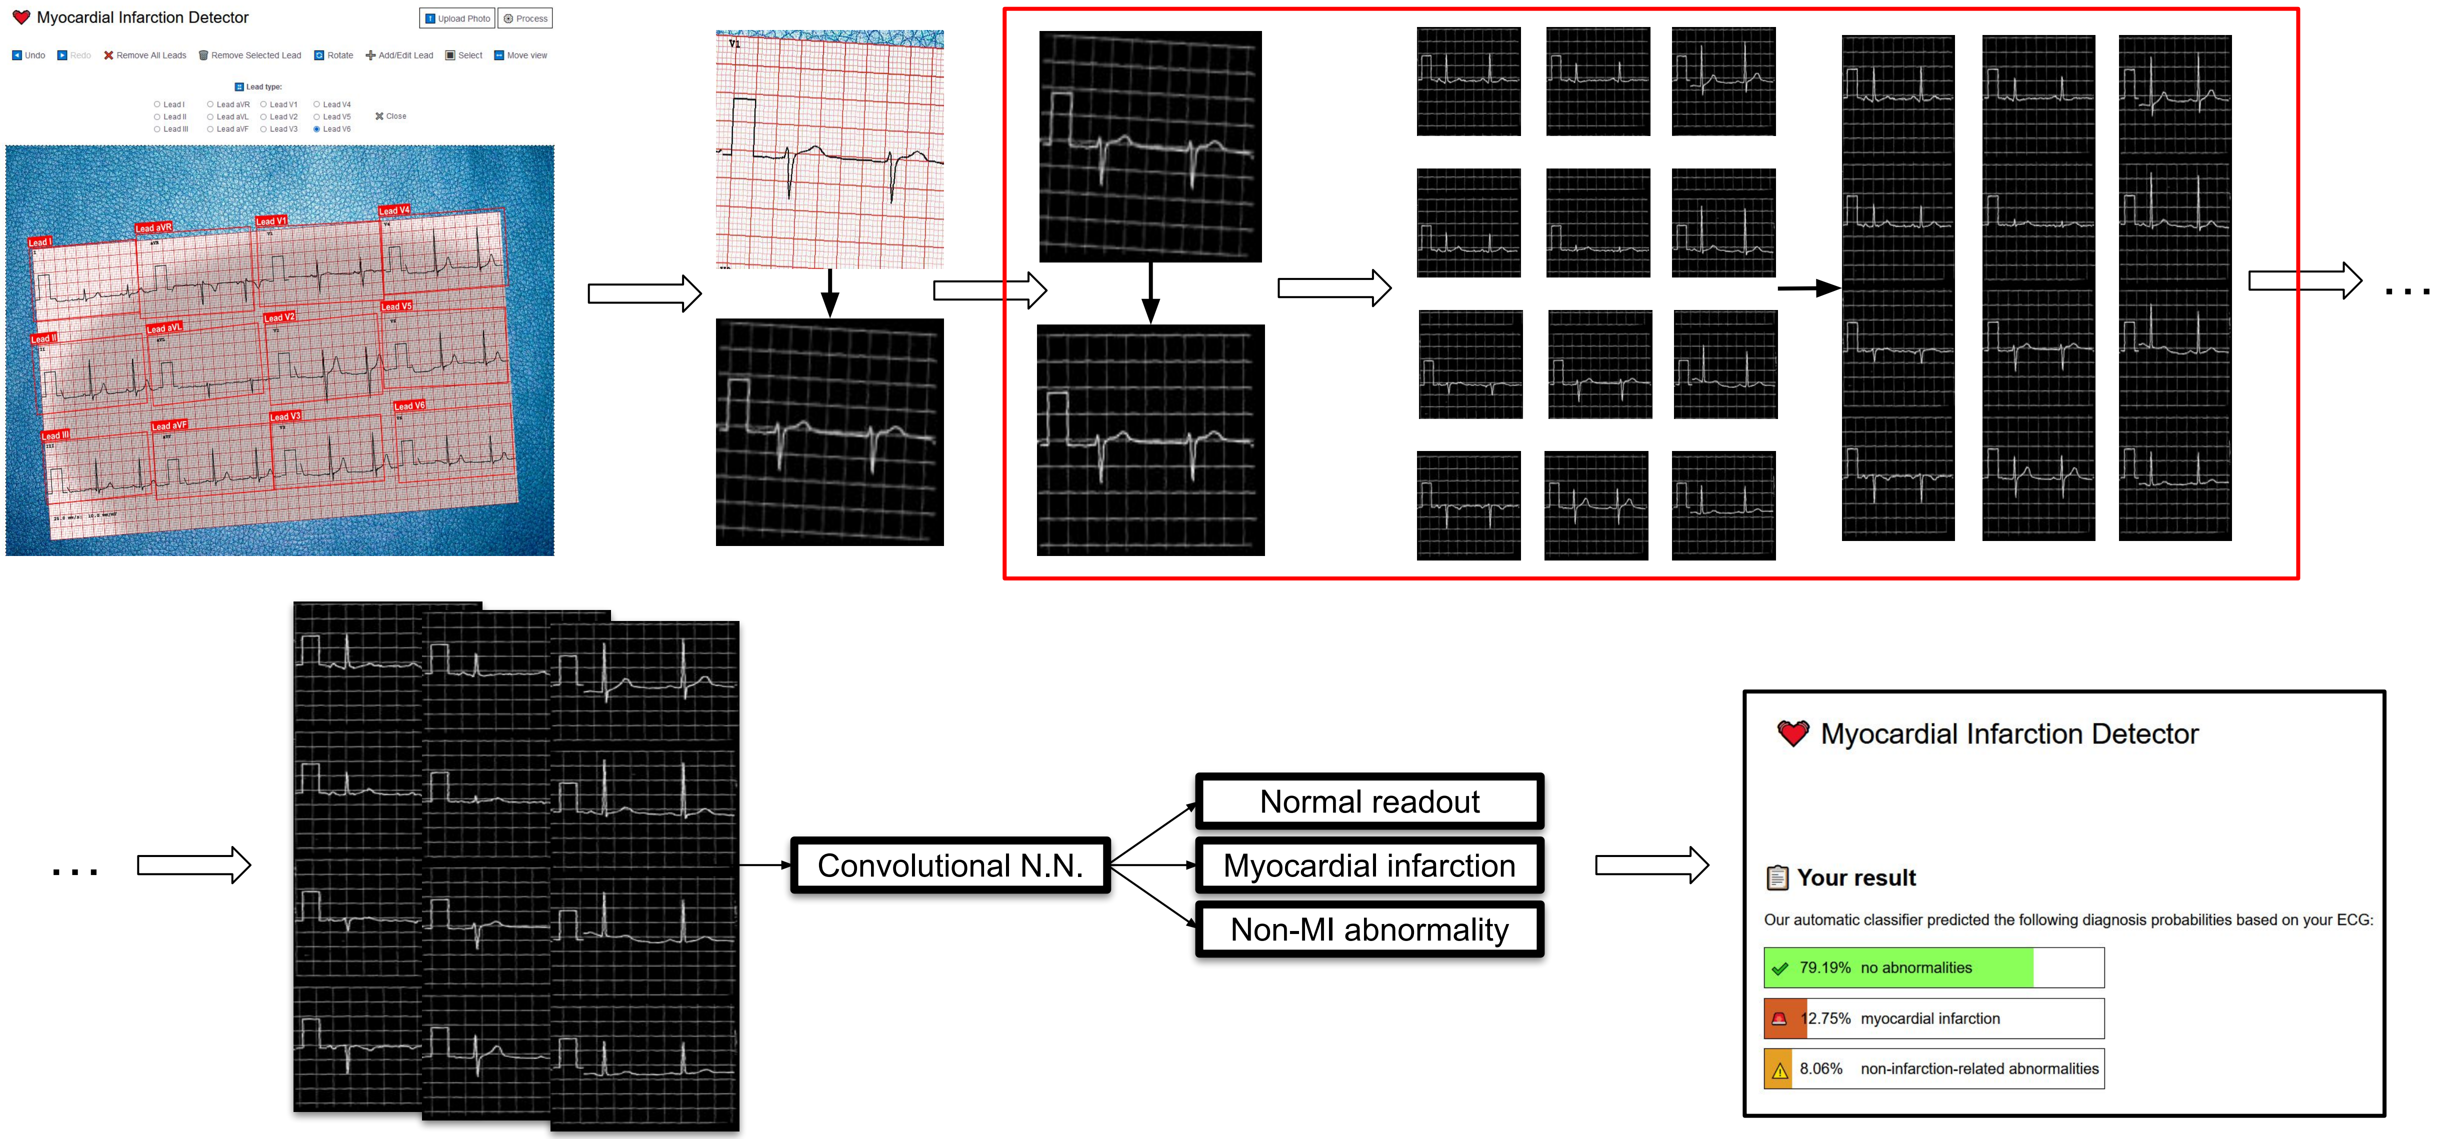

3 ECG Processing Overview

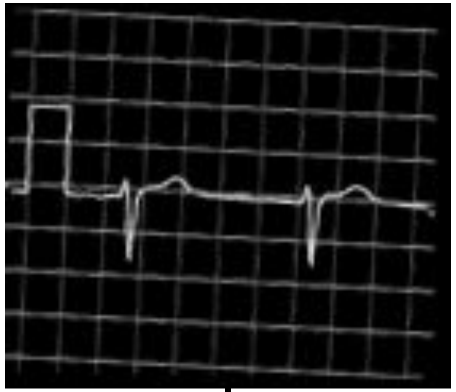

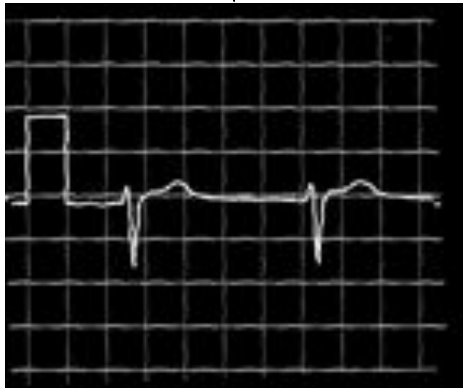

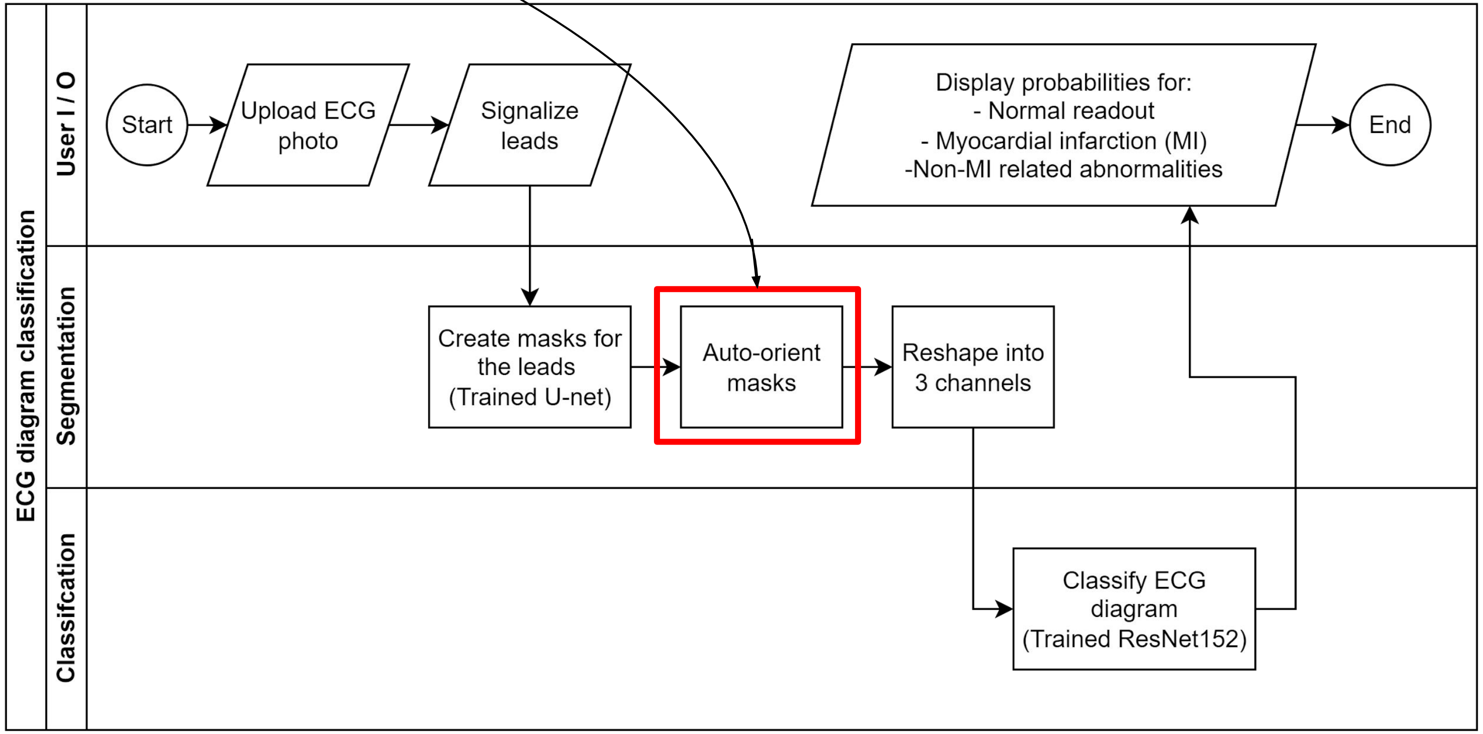

3 ECG Processing Overview

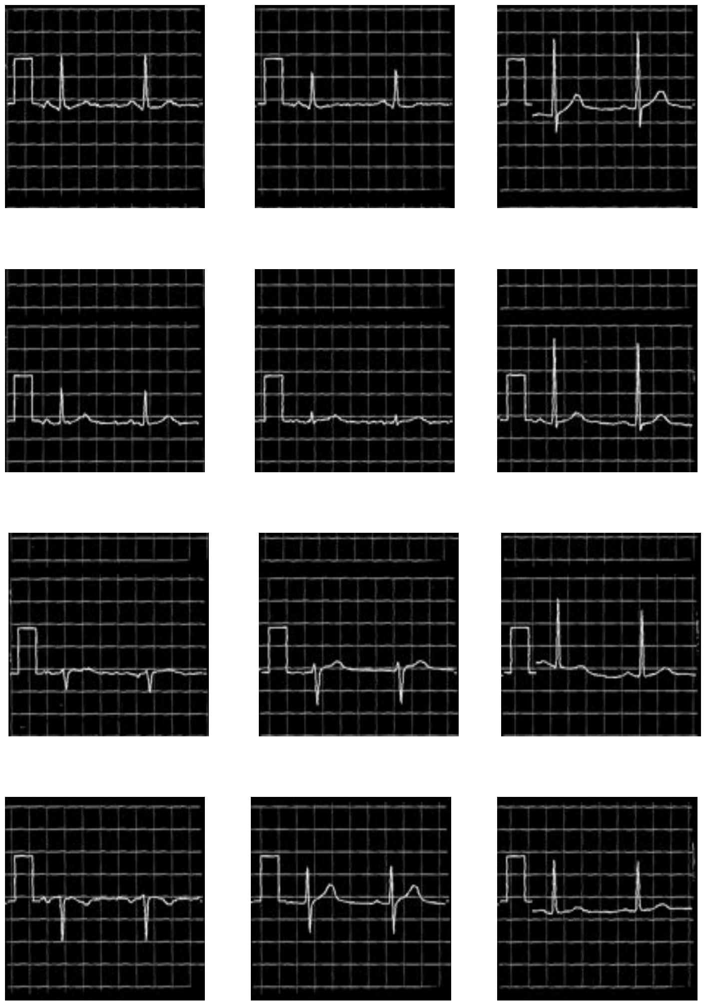

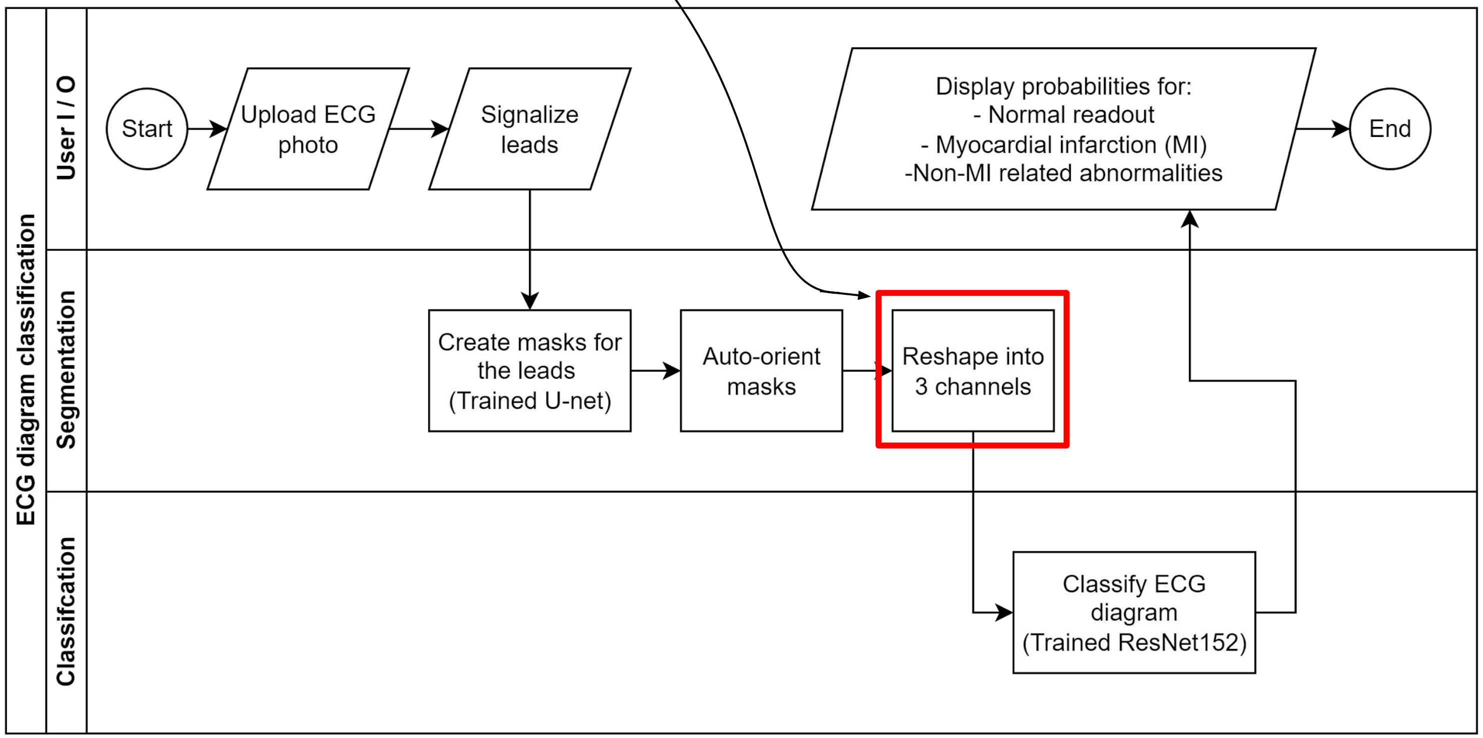

3 ECG Processing Overview

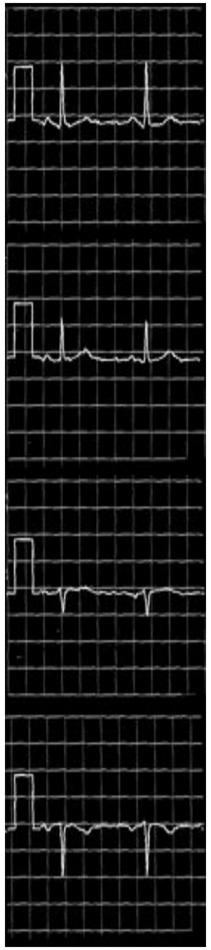

Channel 1

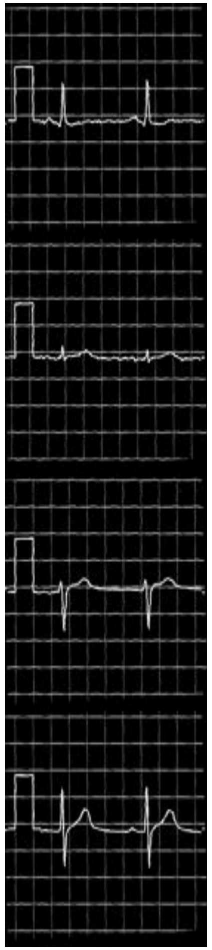

Channel 2

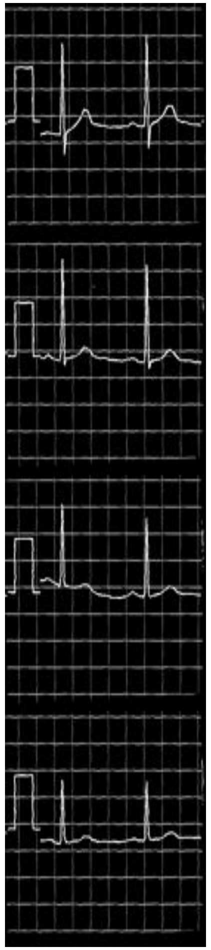

Channel 3

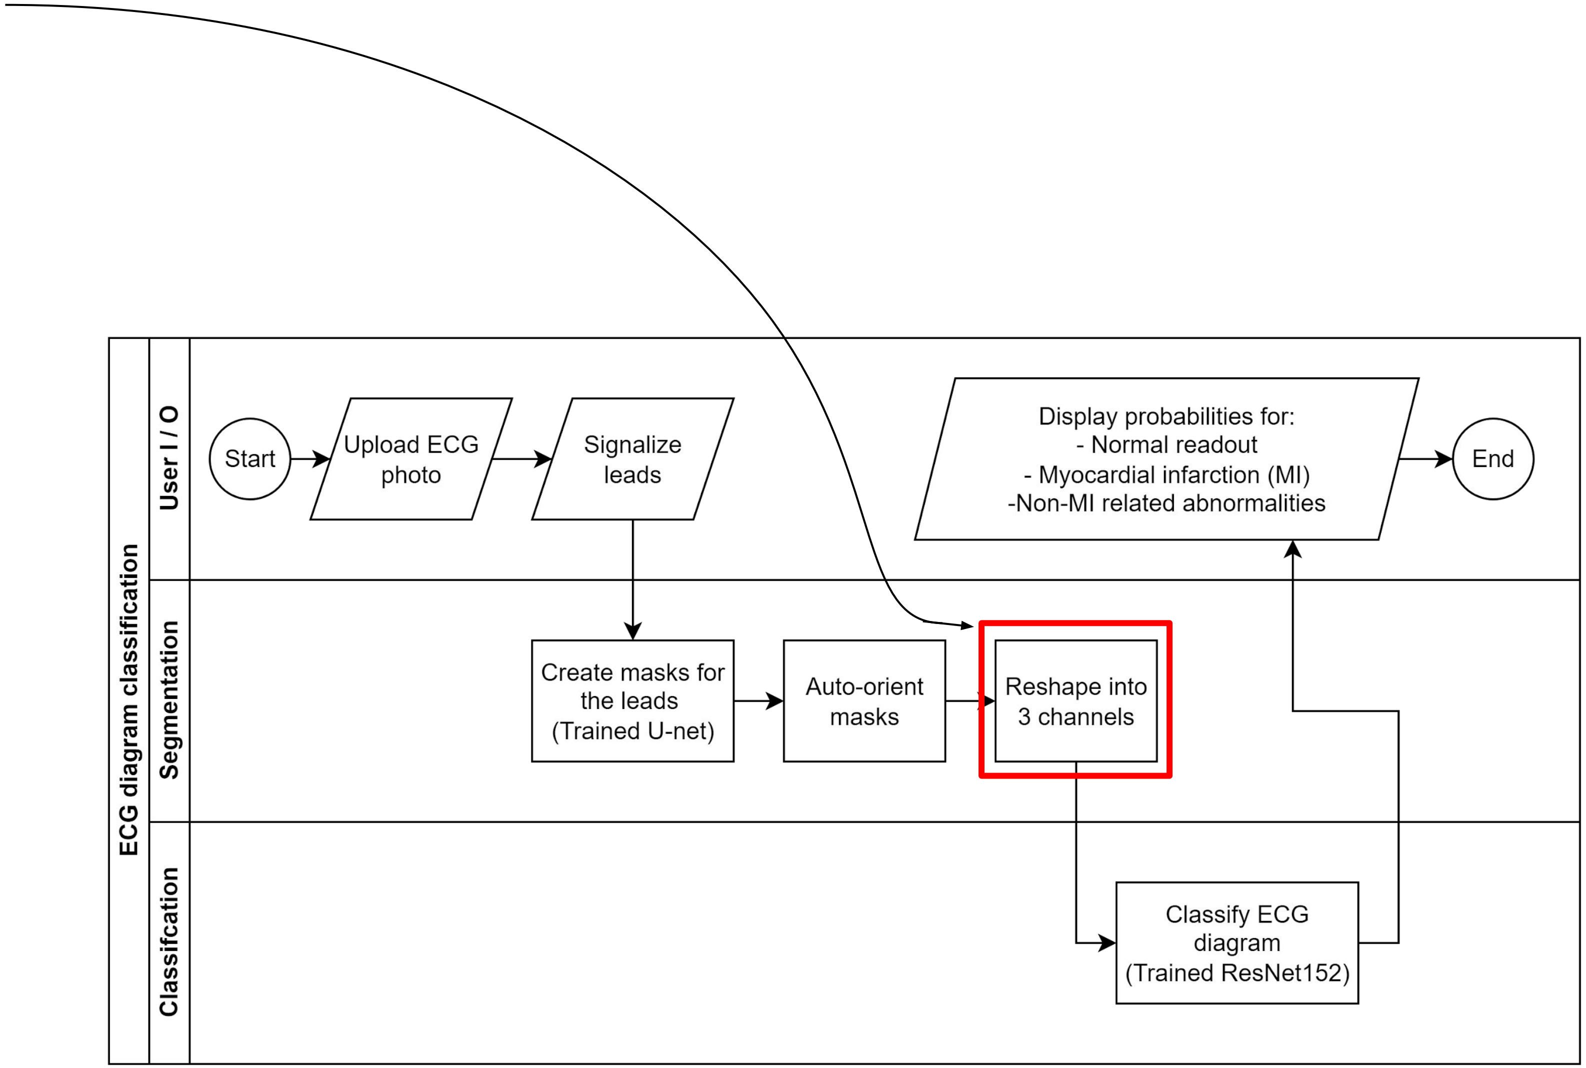

3.1 ECG Segmentation

The straightening algorithm is based on detecting horizontally the grid-lines on several heights and then correct the rotation

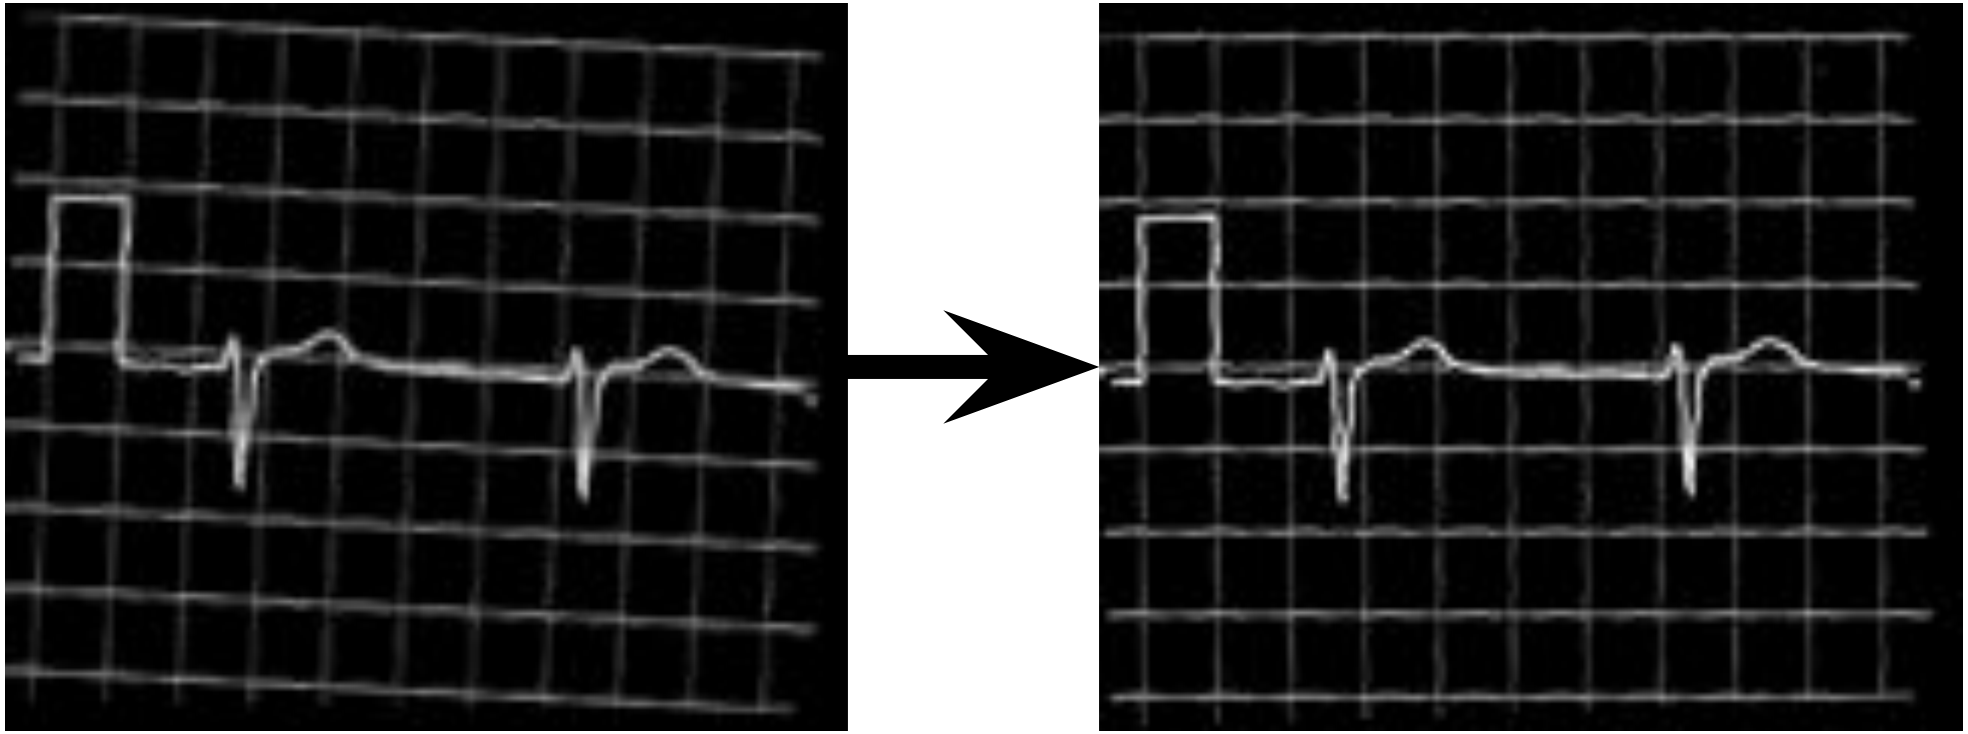

3.1 ECG Segmentation

The 12 leads are divided in 3 groups, one for each channel fed into the convnet for classification.

If more leads are inserted, they are added in a **consistent** way to the different channels. **This is important because certain types of myocardial infarctions result in characteristic changes to certain leads.**

The grid lines are kept also fed to the classifying convnet. We expect that this helps with detecting correlations of different lead values at equal times.

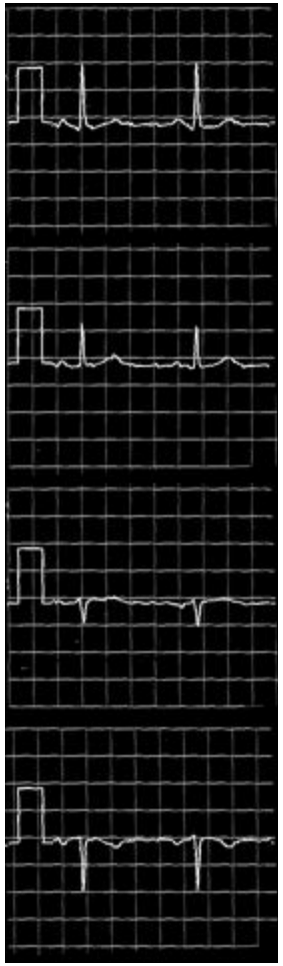

Channel 1

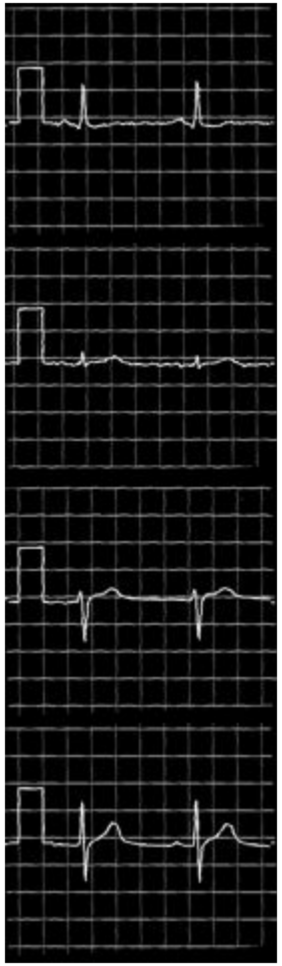

Channel 2

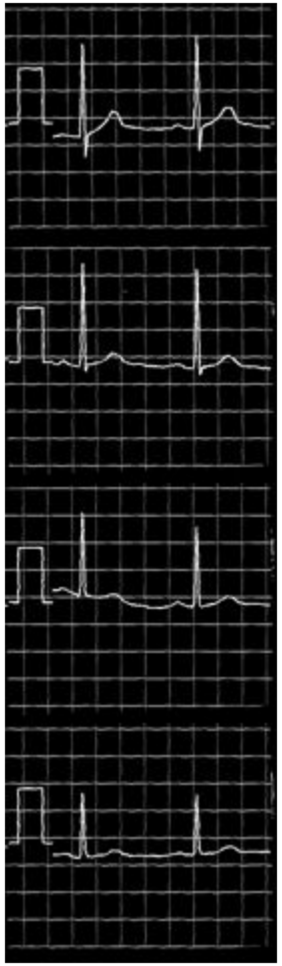

Channel 3

3 ECG Processing Overview

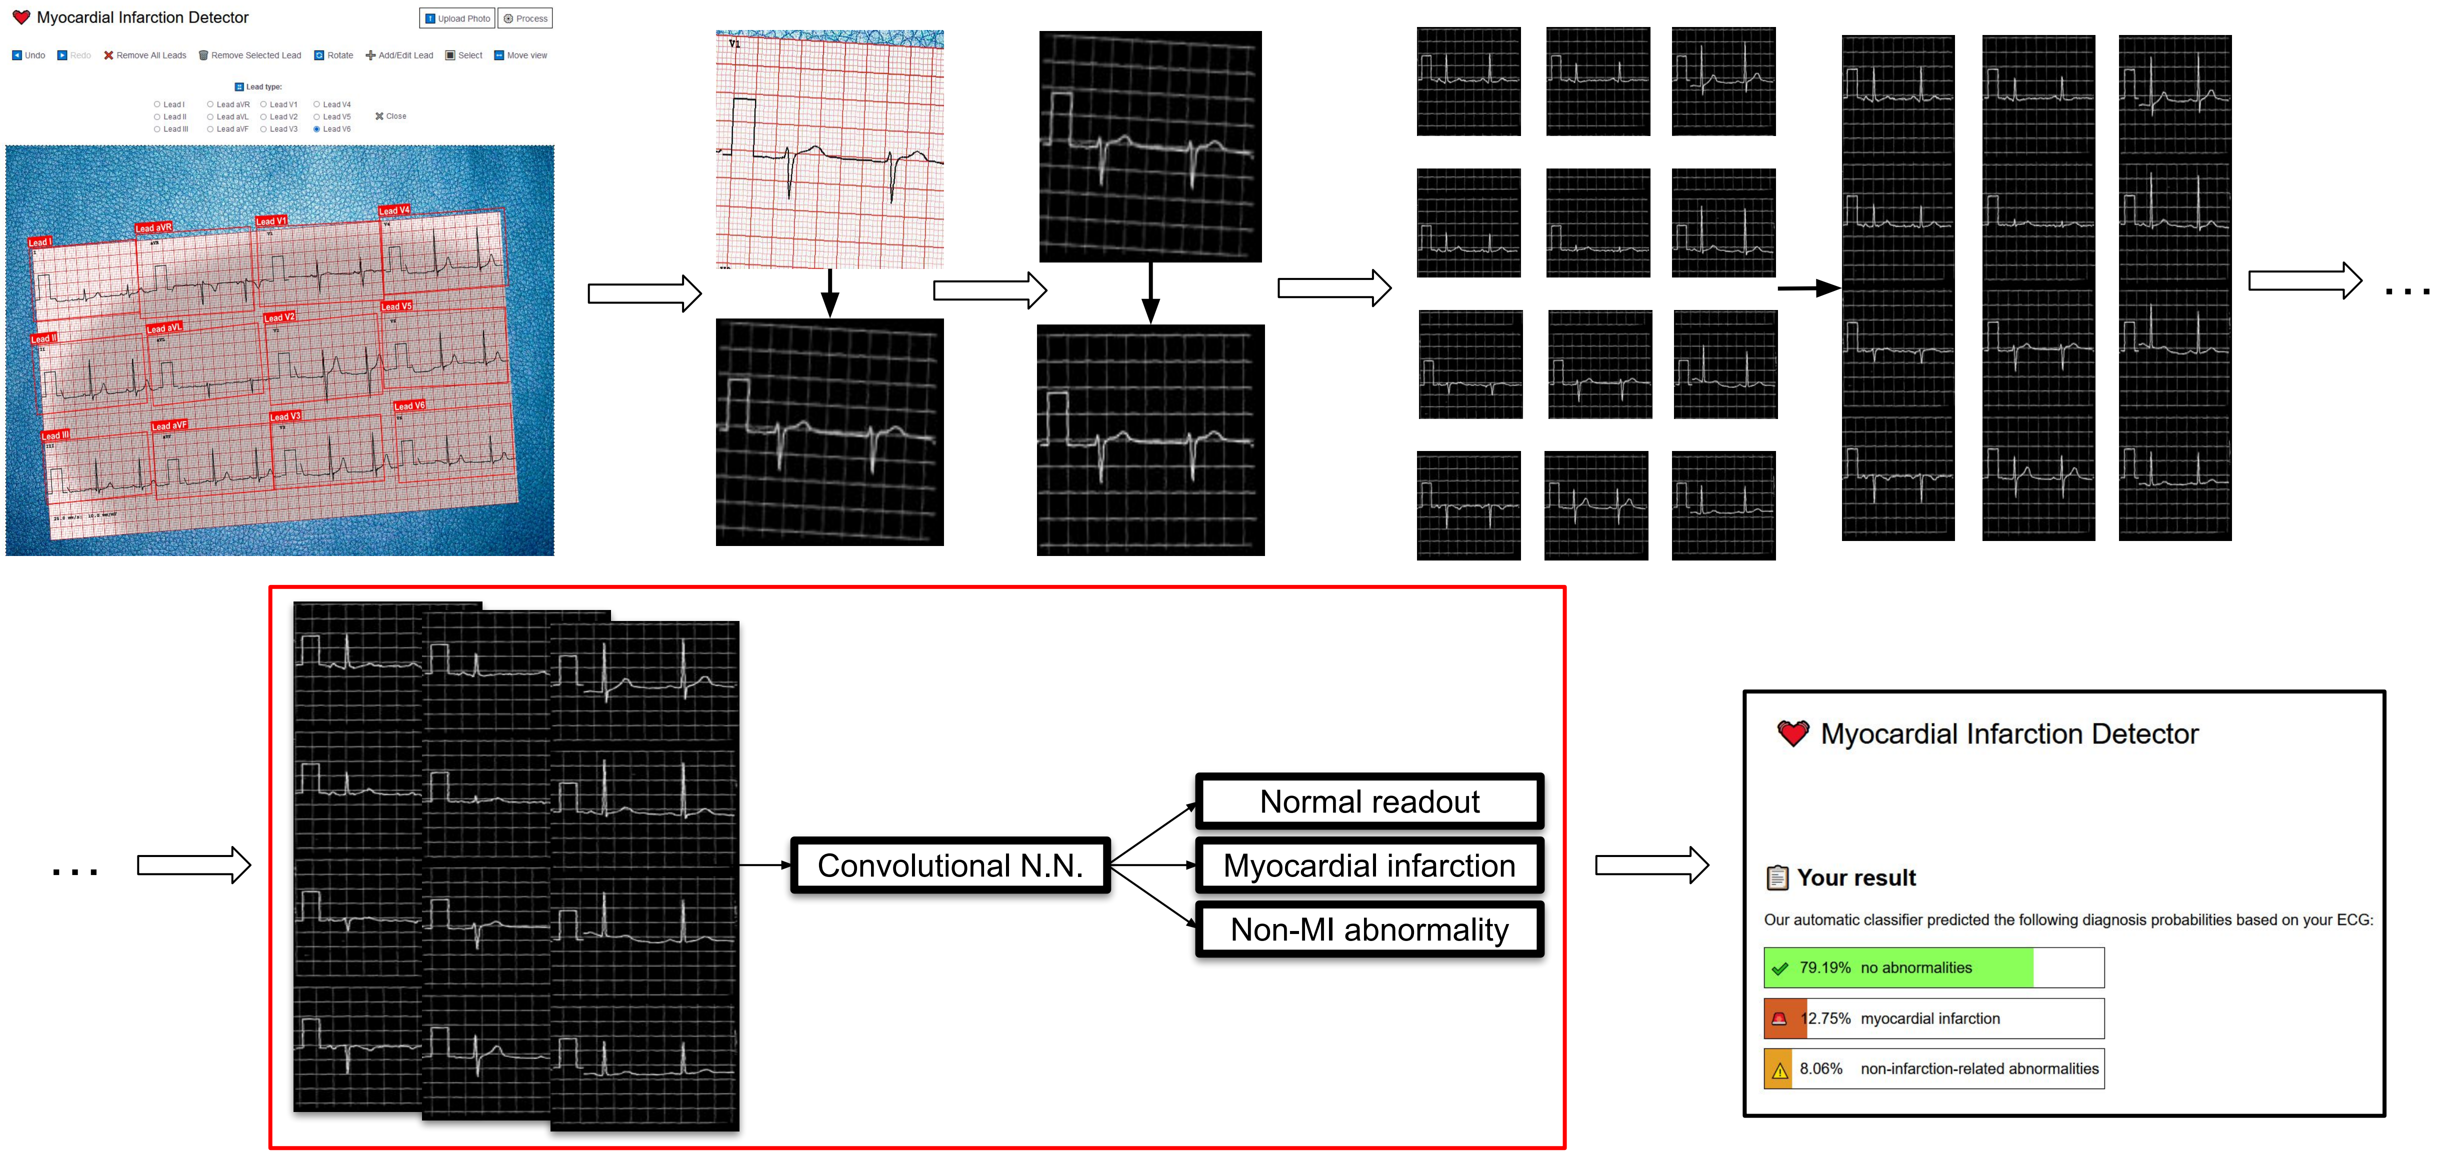

3 ECG Processing Overview

Channel 1 Channel 2 Channel 3

ResNet152

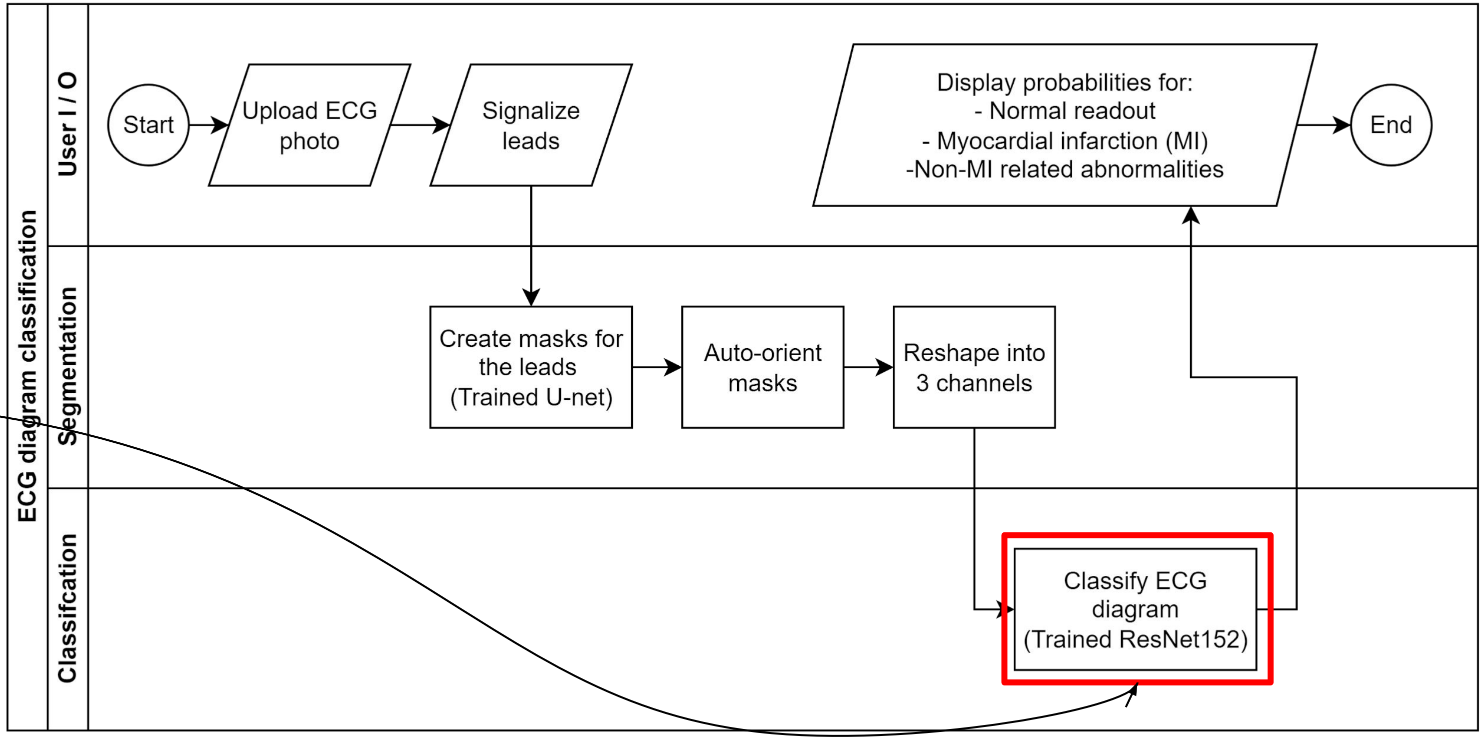

3.2 ECG Classification

Channel 1 Channel 2 Channel 3

3.2 ECG Classification

Channel 1 Channel 2 Channel 3

ResNet152

3.2 ECG Classification

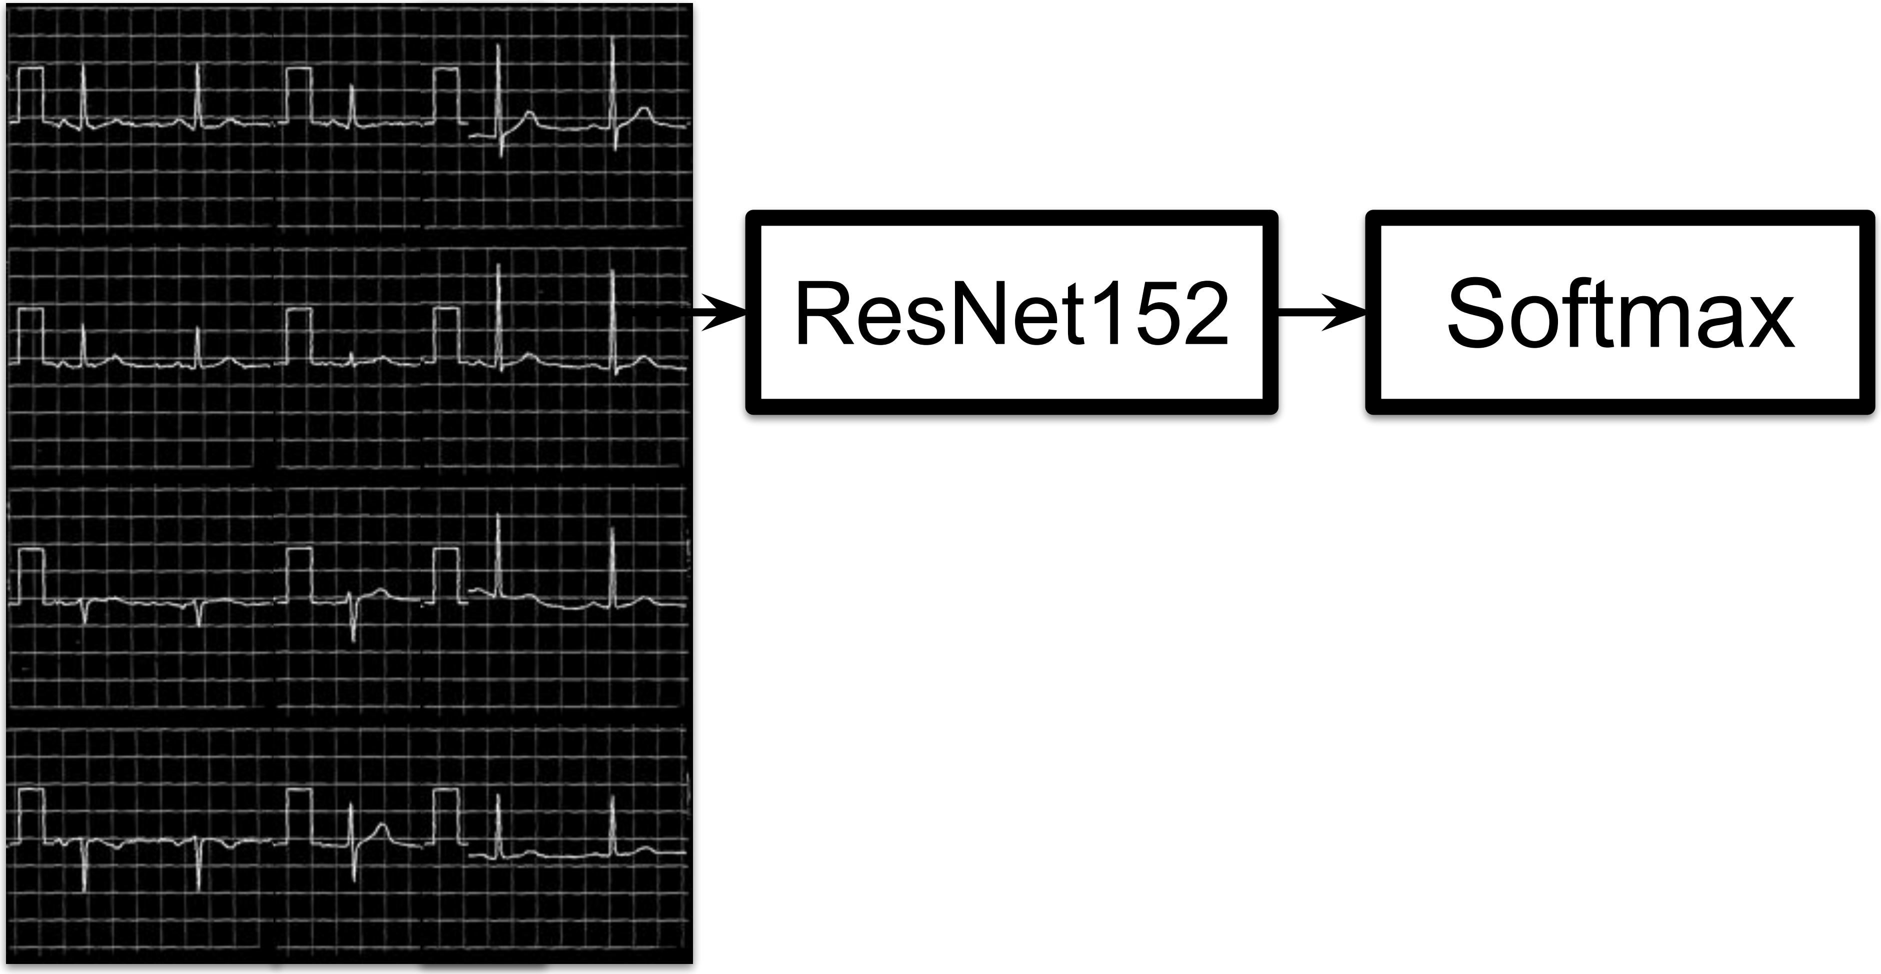

Channel 1 Channel 2 Channel 3

3.2 ECG Classification

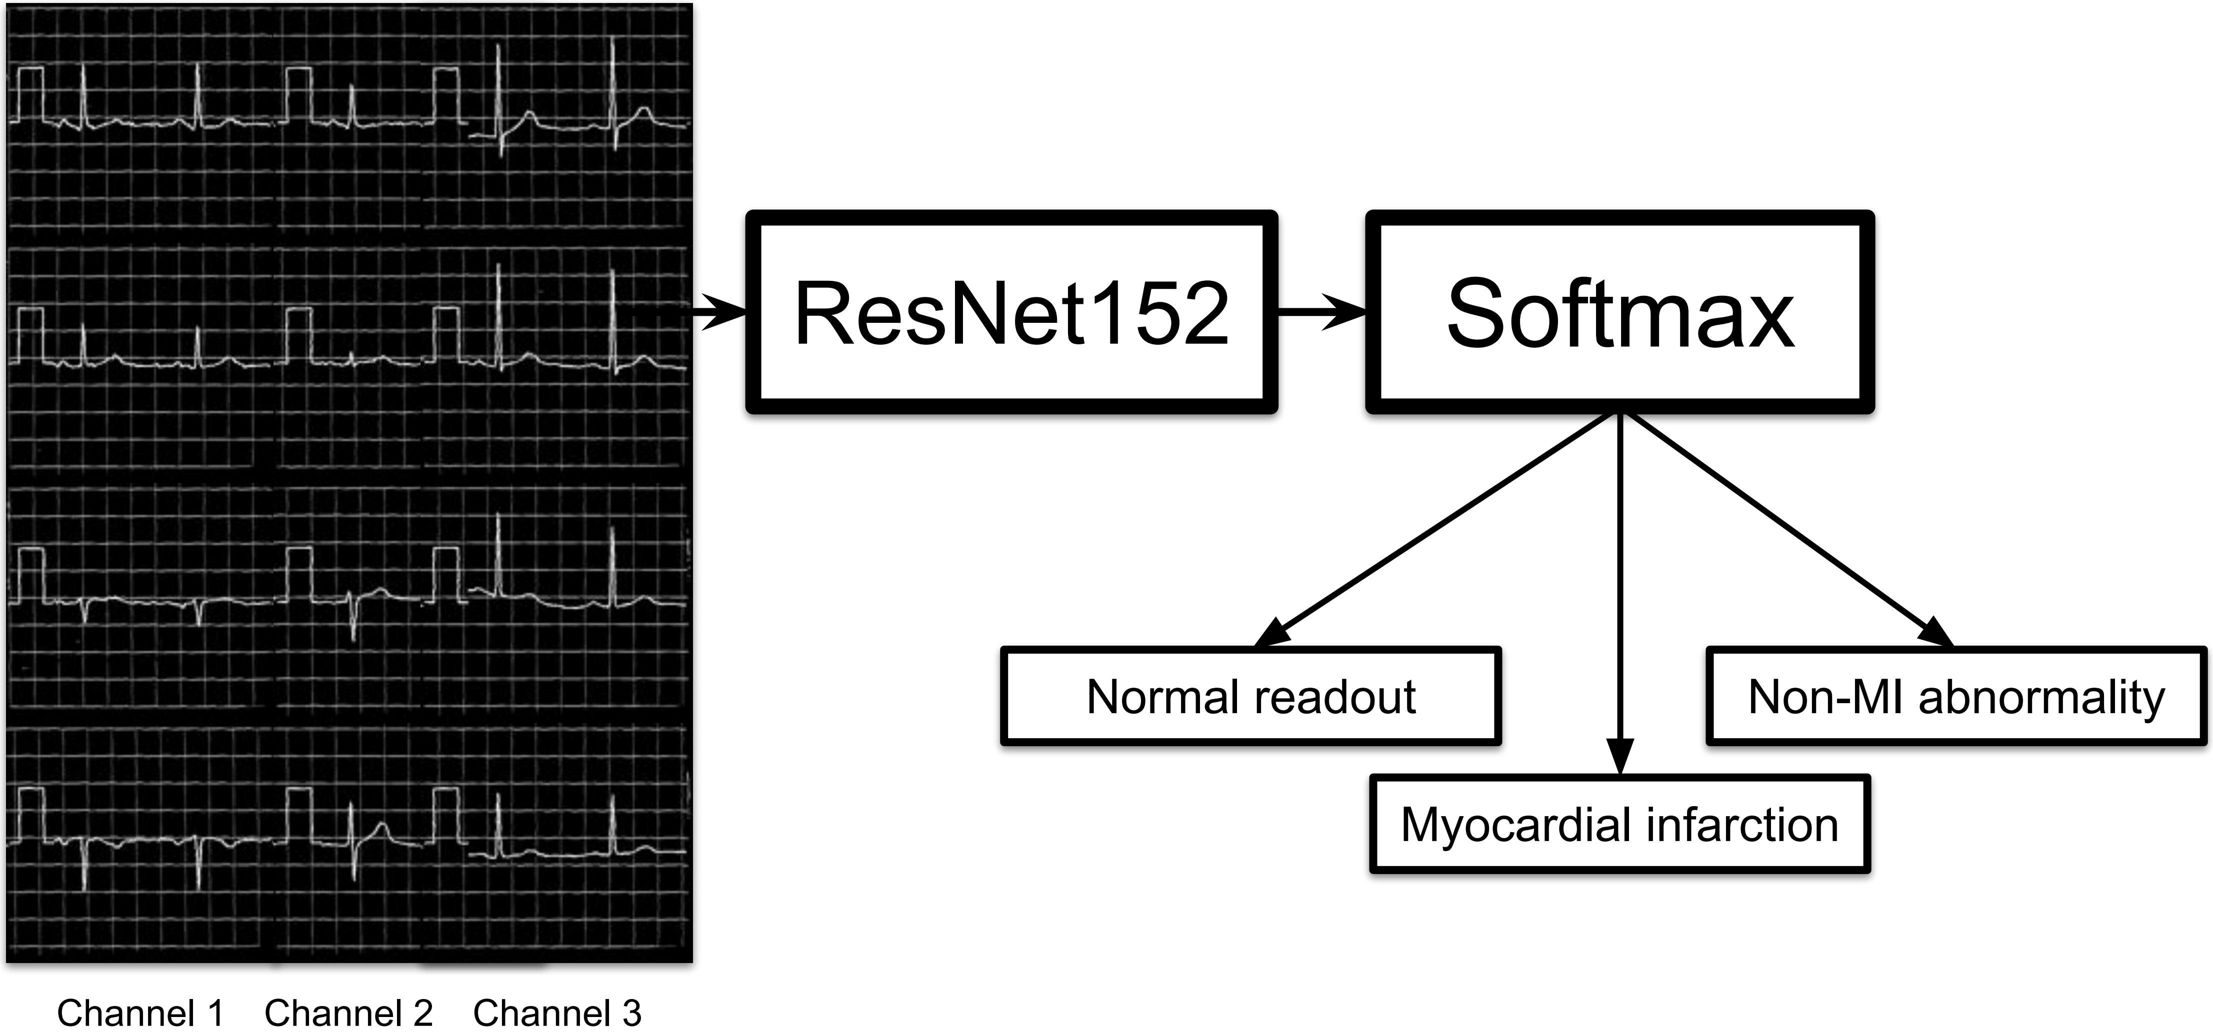

3.2 ECG Classification: ResNet152 architecture

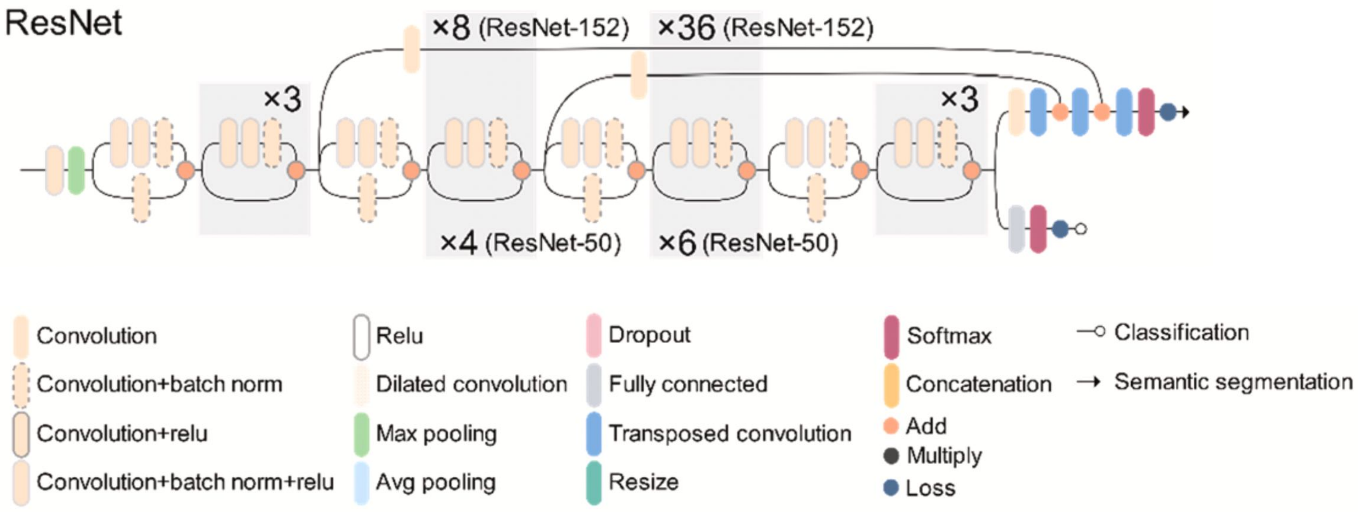

Figure 1: ResNet152 architecture diagram

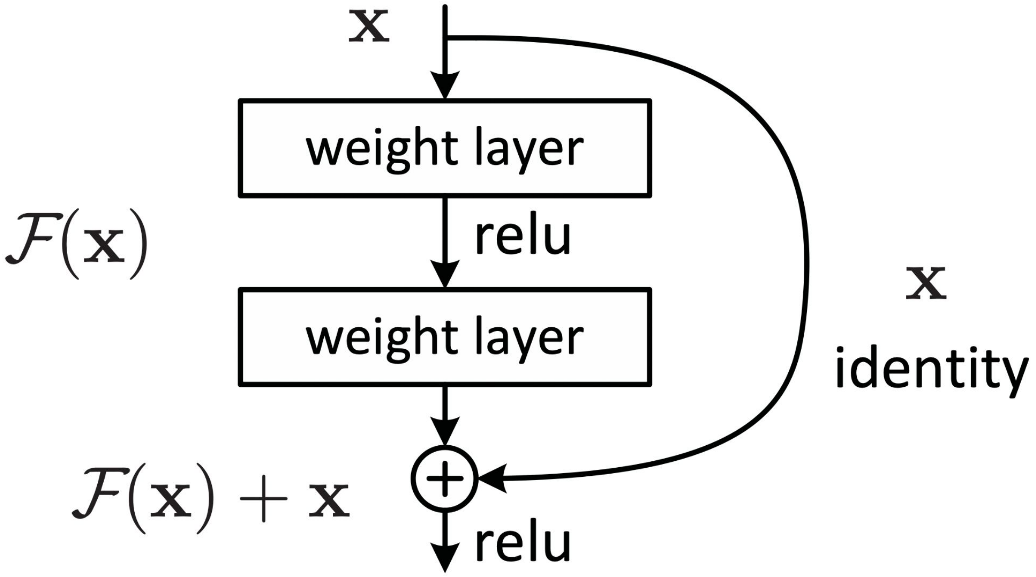

Figure 2: Residual learning building block

Figure 1 taken from: Zhouxin Xi, Chris Hopkinson, Stewart B. Rood, & Derek R. Peddle (2020). See the forest and the trees: Effective machine and deep learning algorithms for wood filtering and tree species classification from terrestrial laser scanning. ISPRS Journal of Photogrammetry and Remote Sensing, 168, 1-16.

Figure 2 taken from: Kaiming He, Xiangyu Zhang, Shaoqing Ren, & Jian Sun (2015). Deep Residual Learning for Image Recognition. CoRR, https://arxiv.org/abs/1512.03385v1

3 ECG Processing Overview

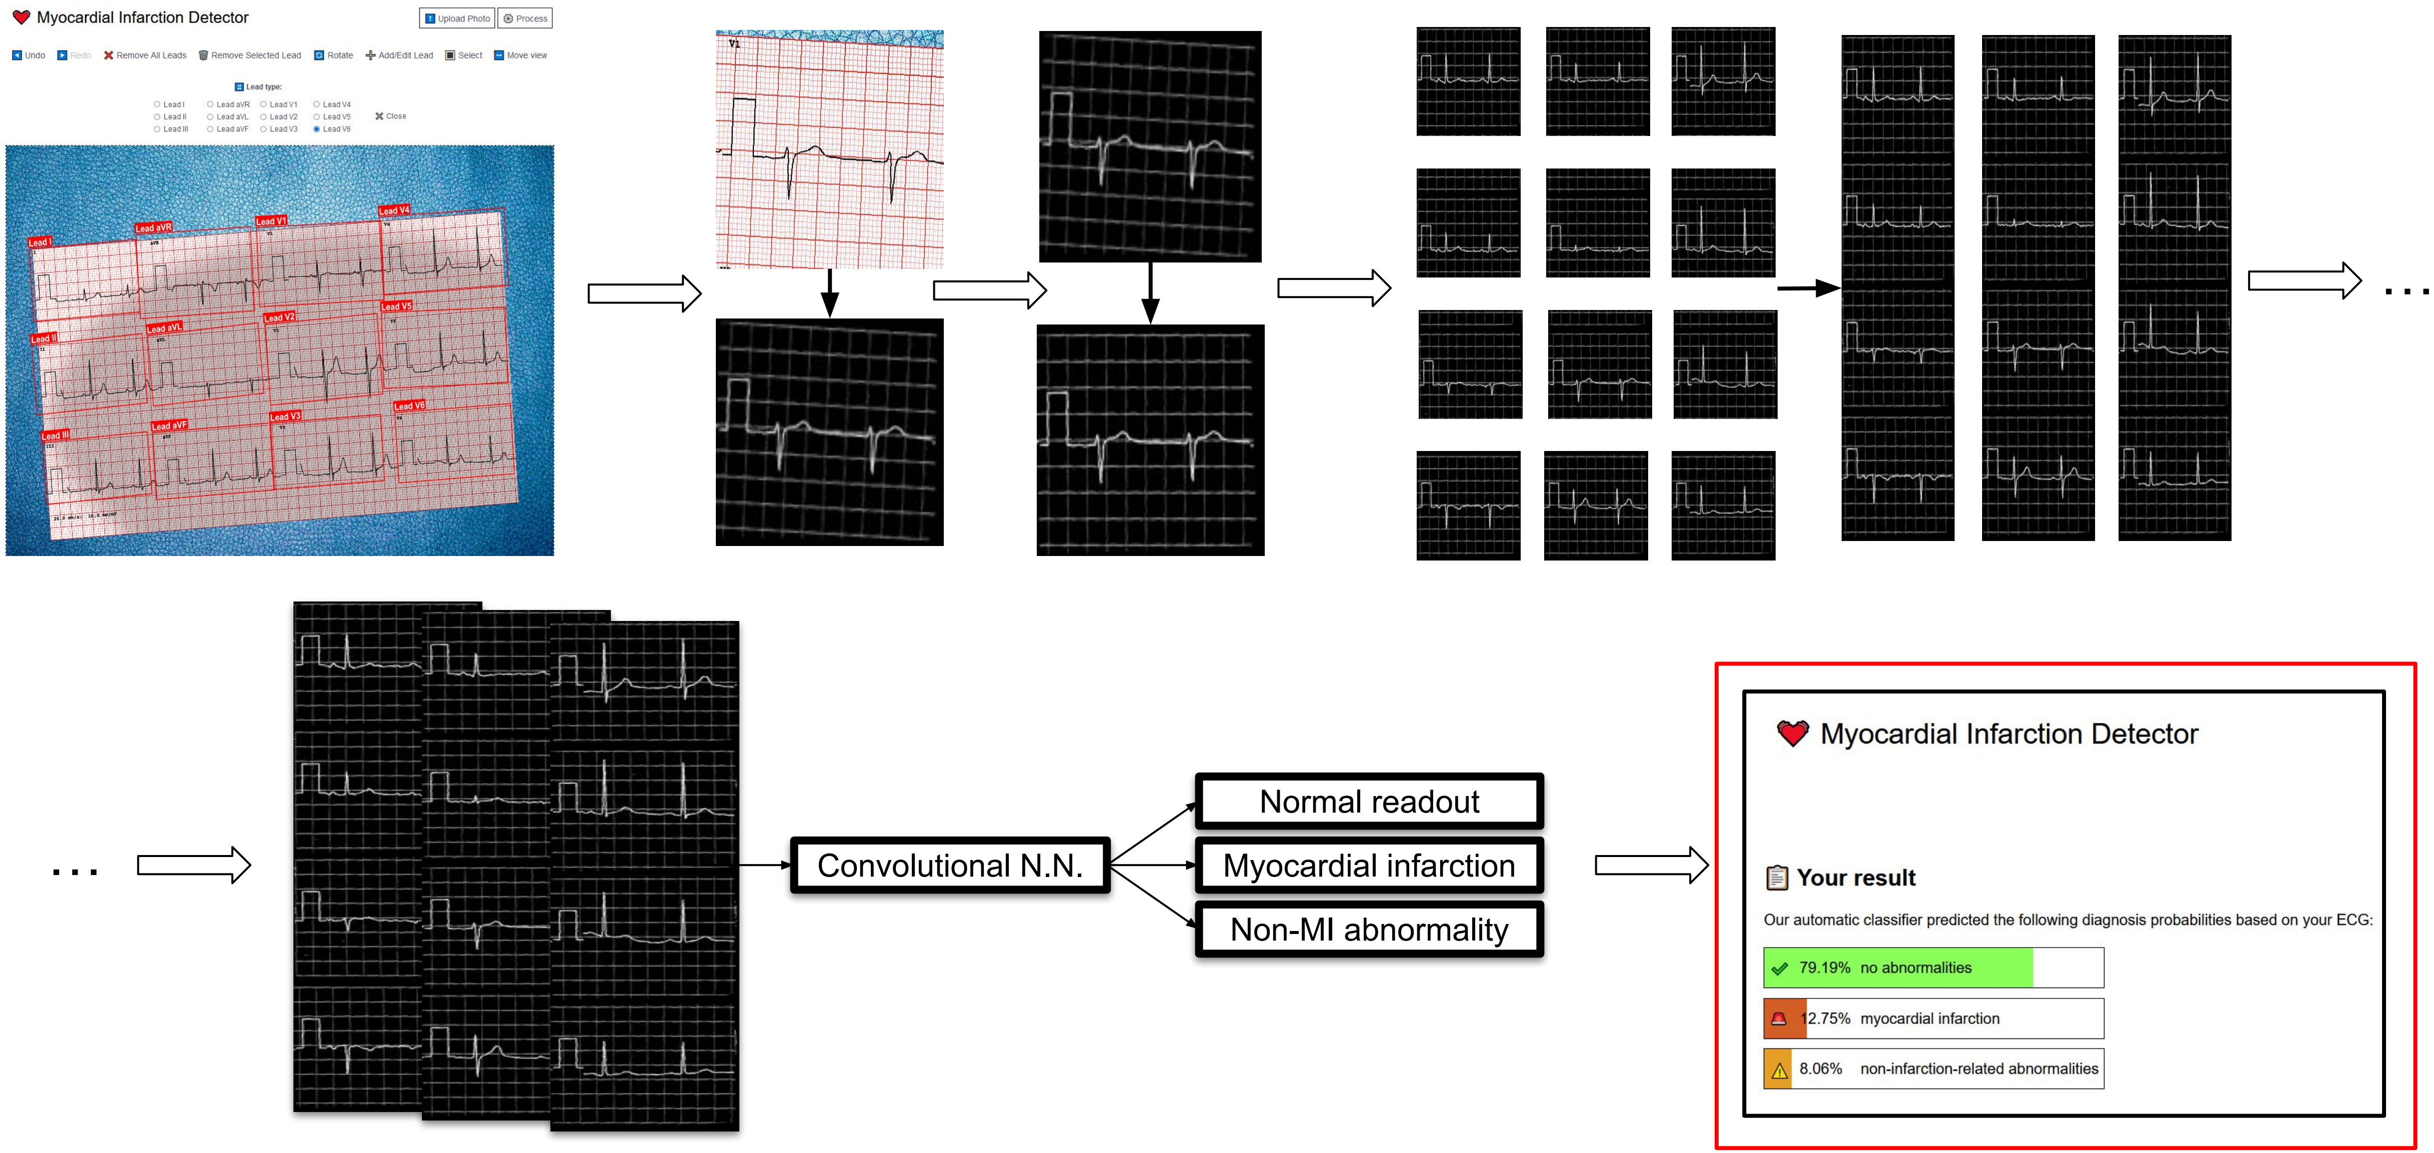

3 ECG Processing Overview

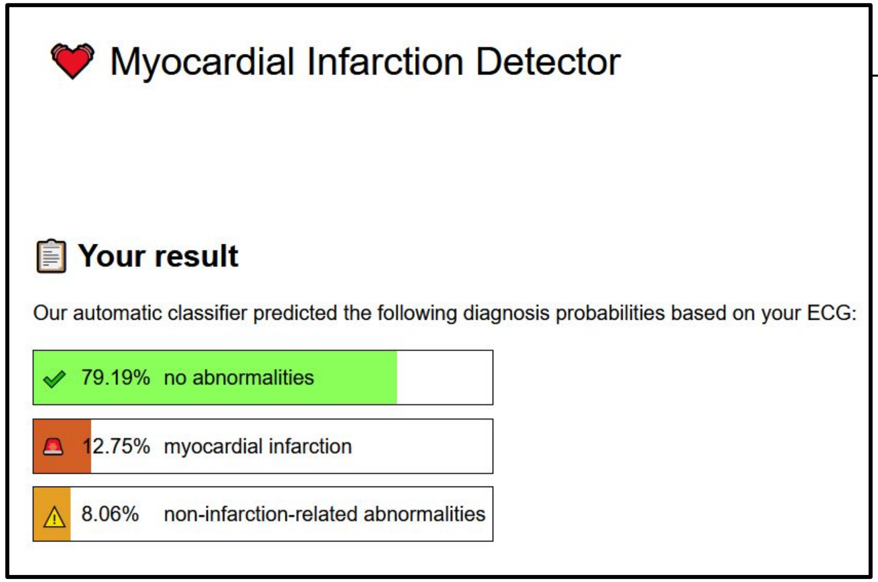

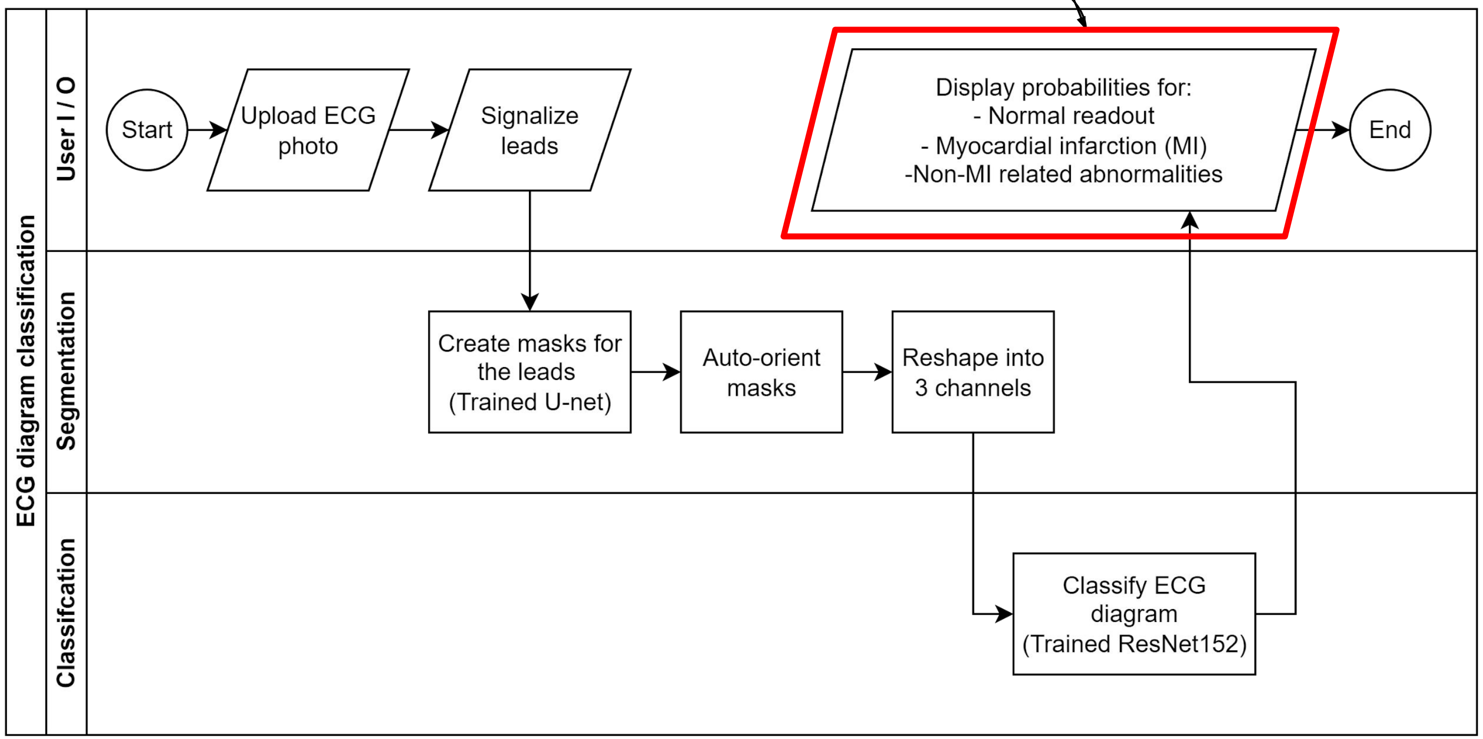

This will be presented with detail later on.

3.3 ECG Network Training

Extraction of segmentation images

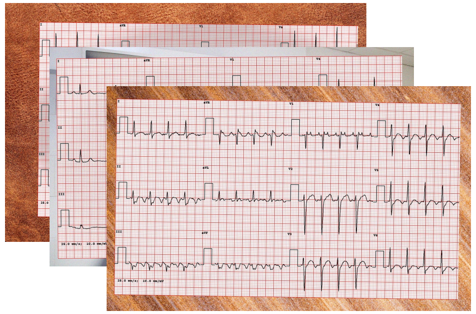

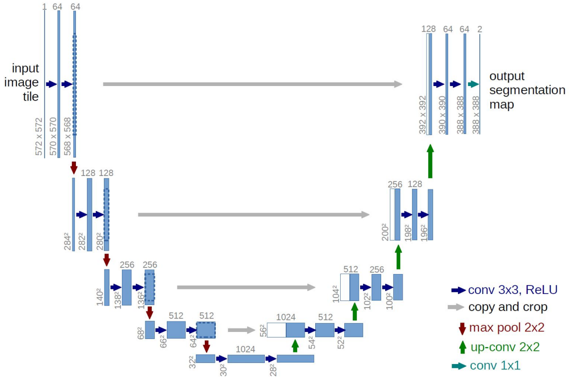

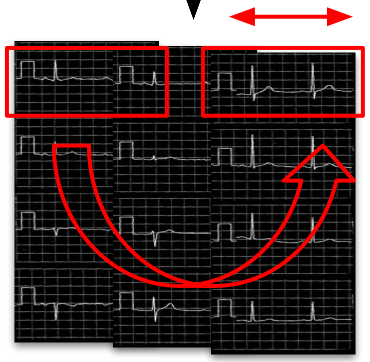

Swapping and stretching of leads

Data augmentation

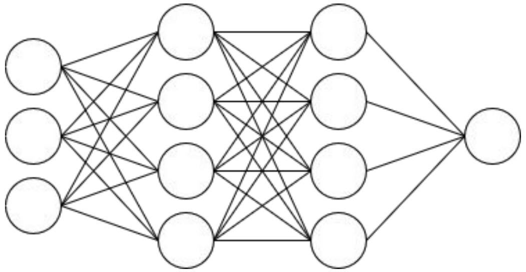

Training on augmented segmentation images

dataset stats:

before augmentation: 13826 ECGs
6152 normal ECGs
4752 ECGs with non-MI abnormalities
2922 ECGs showing MIs

after augmentation: 41478 ECGs
18456 normal ECGs
14256 ECGs with non-MI abnormalities
8766 ECGs showing MIs

3.3 ECG Network Training: Training Environment

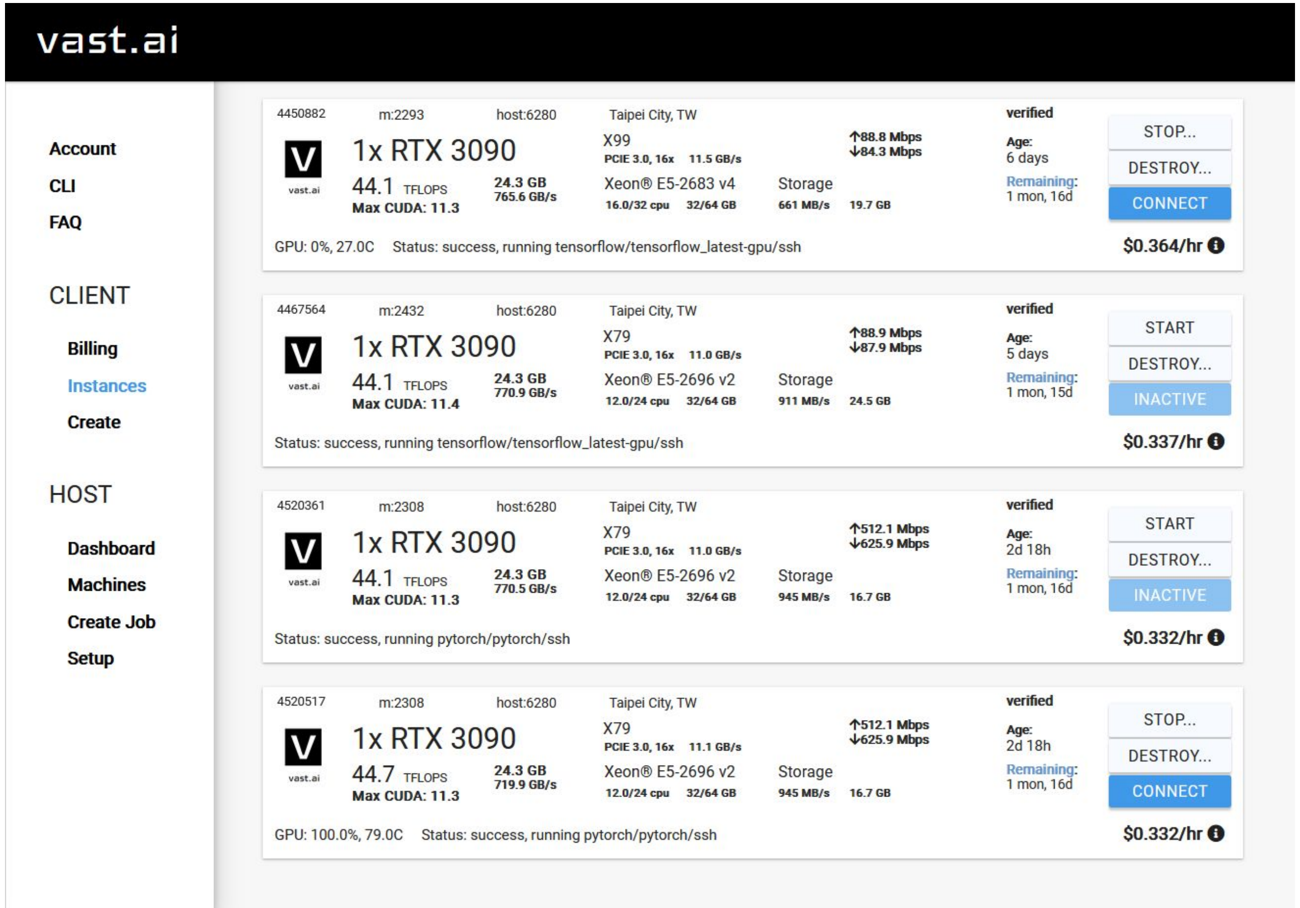

Remote development & training using vast.ai

3.3 ECG Network Training: Training Frameworks

Frameworks:

For classification network

For segmentation network

3.3 ECG Network Training: Training Frameworks

Frameworks:

fast.ai

mlflow

TensorFlow

For classification network

For segmentation network

3.3 ECG Network Training: Training Frameworks

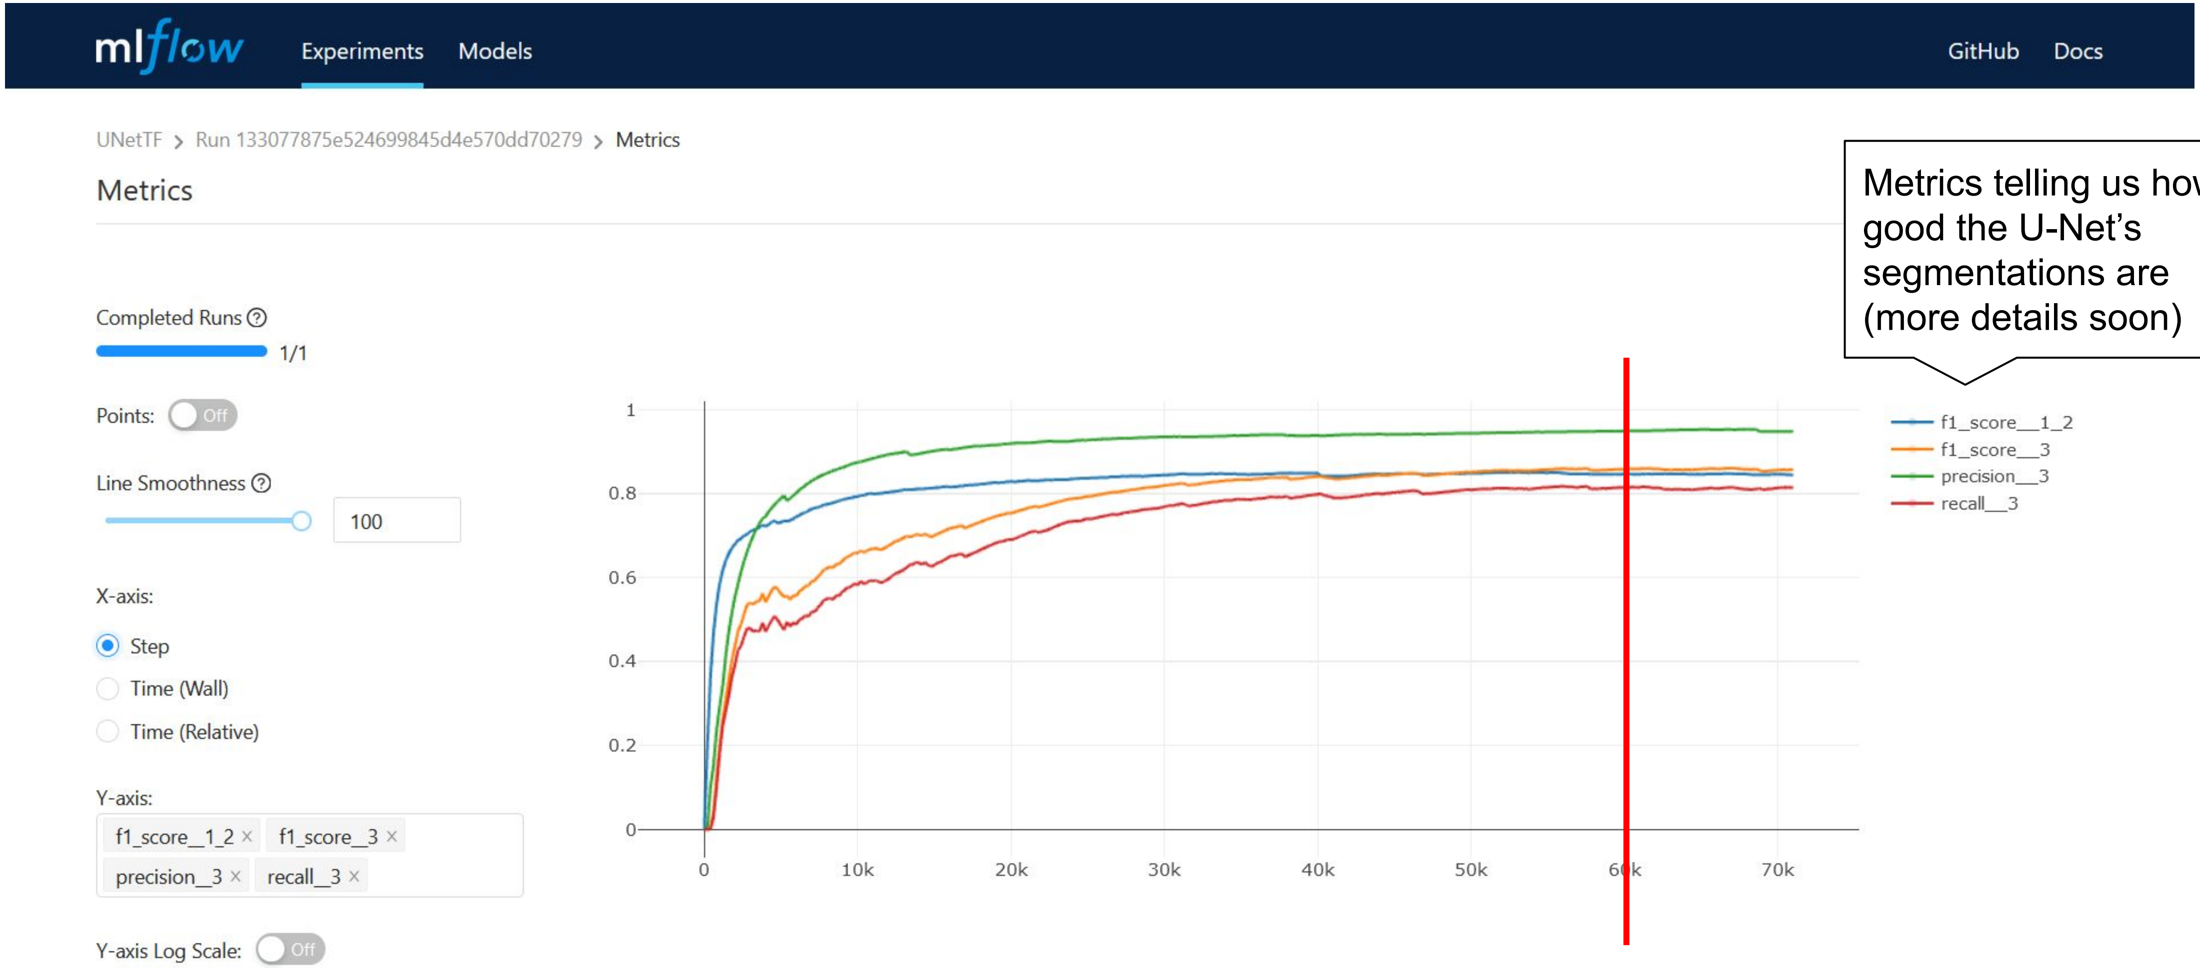

Metrics telling us how good the U-Net's segmentations are (more details soon)

We can easily see (approximate) the step at which convergence is achieved.

3.3 ECG Network Training: Training Frameworks

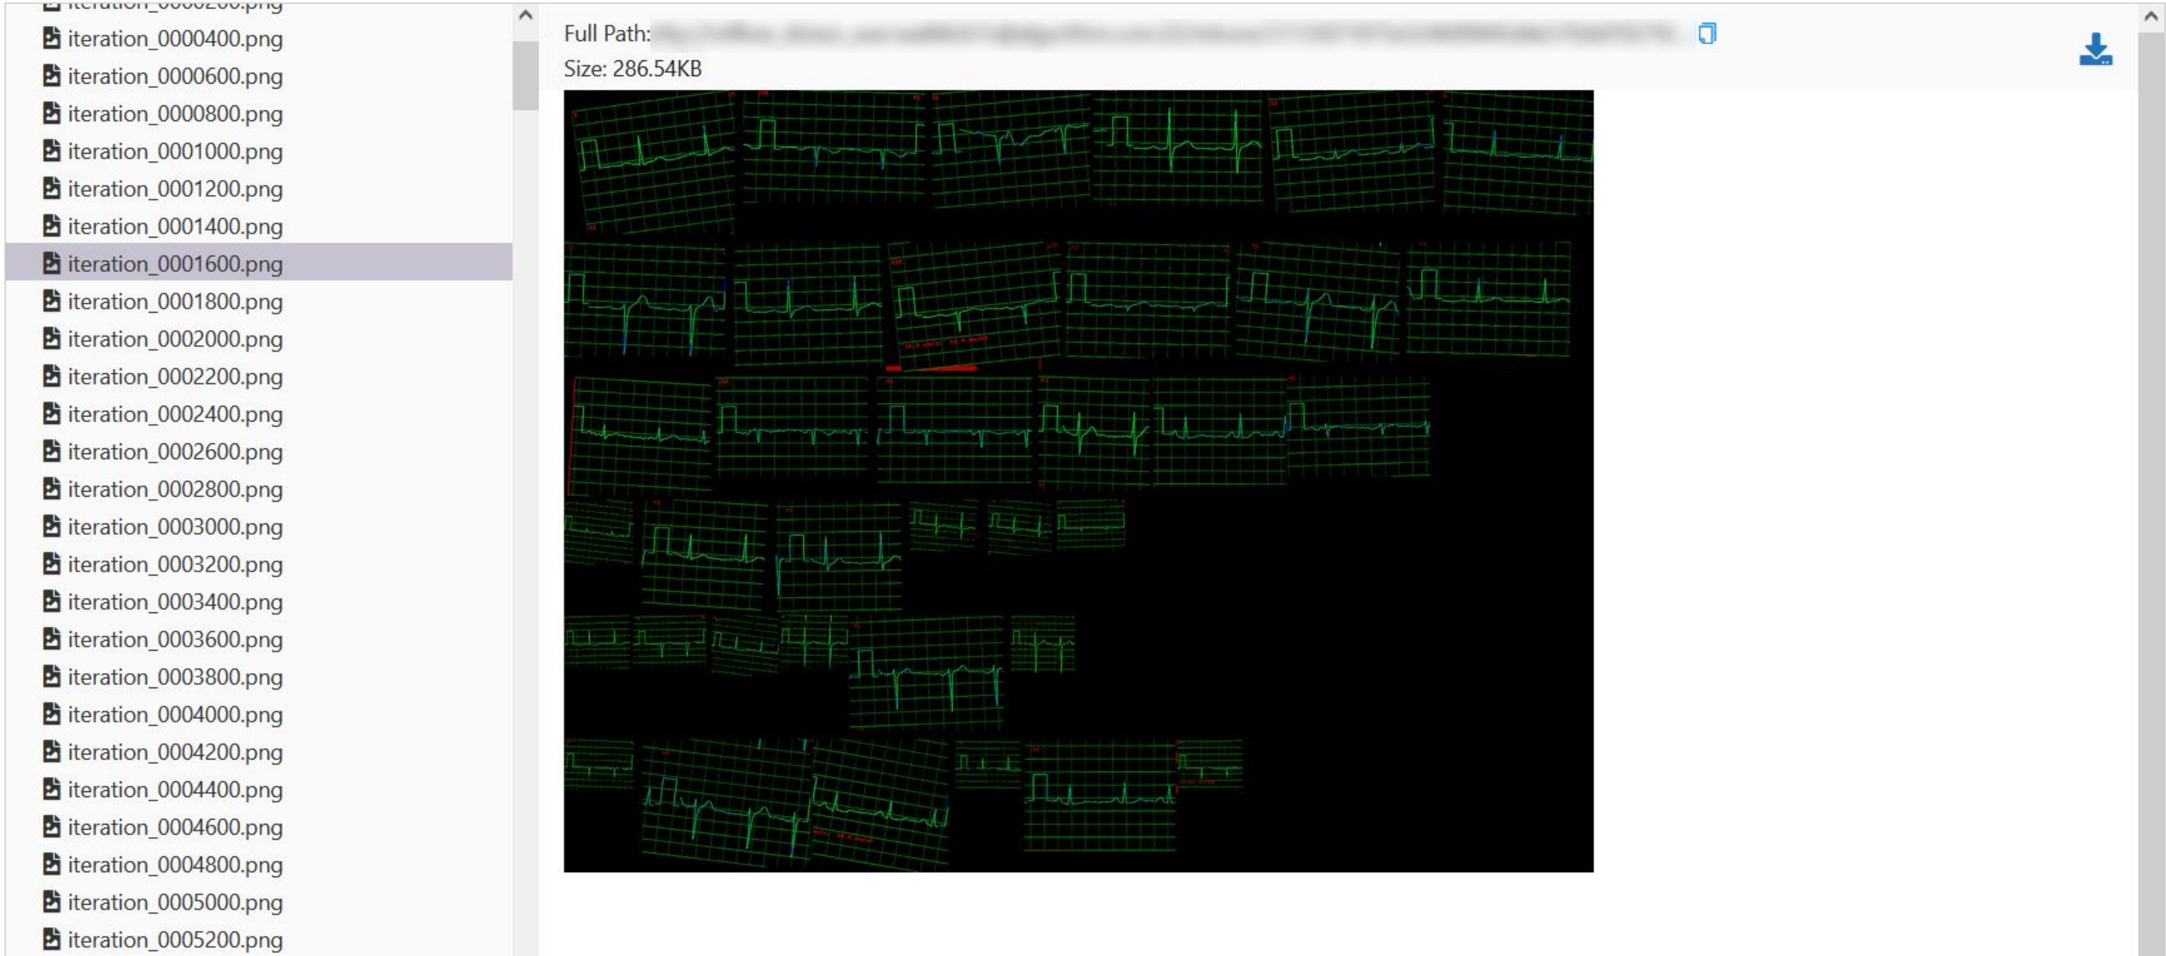

We can easily inspect validation set segmentations at a particular step of training.

3.3 ECG Network Training: Training Frameworks

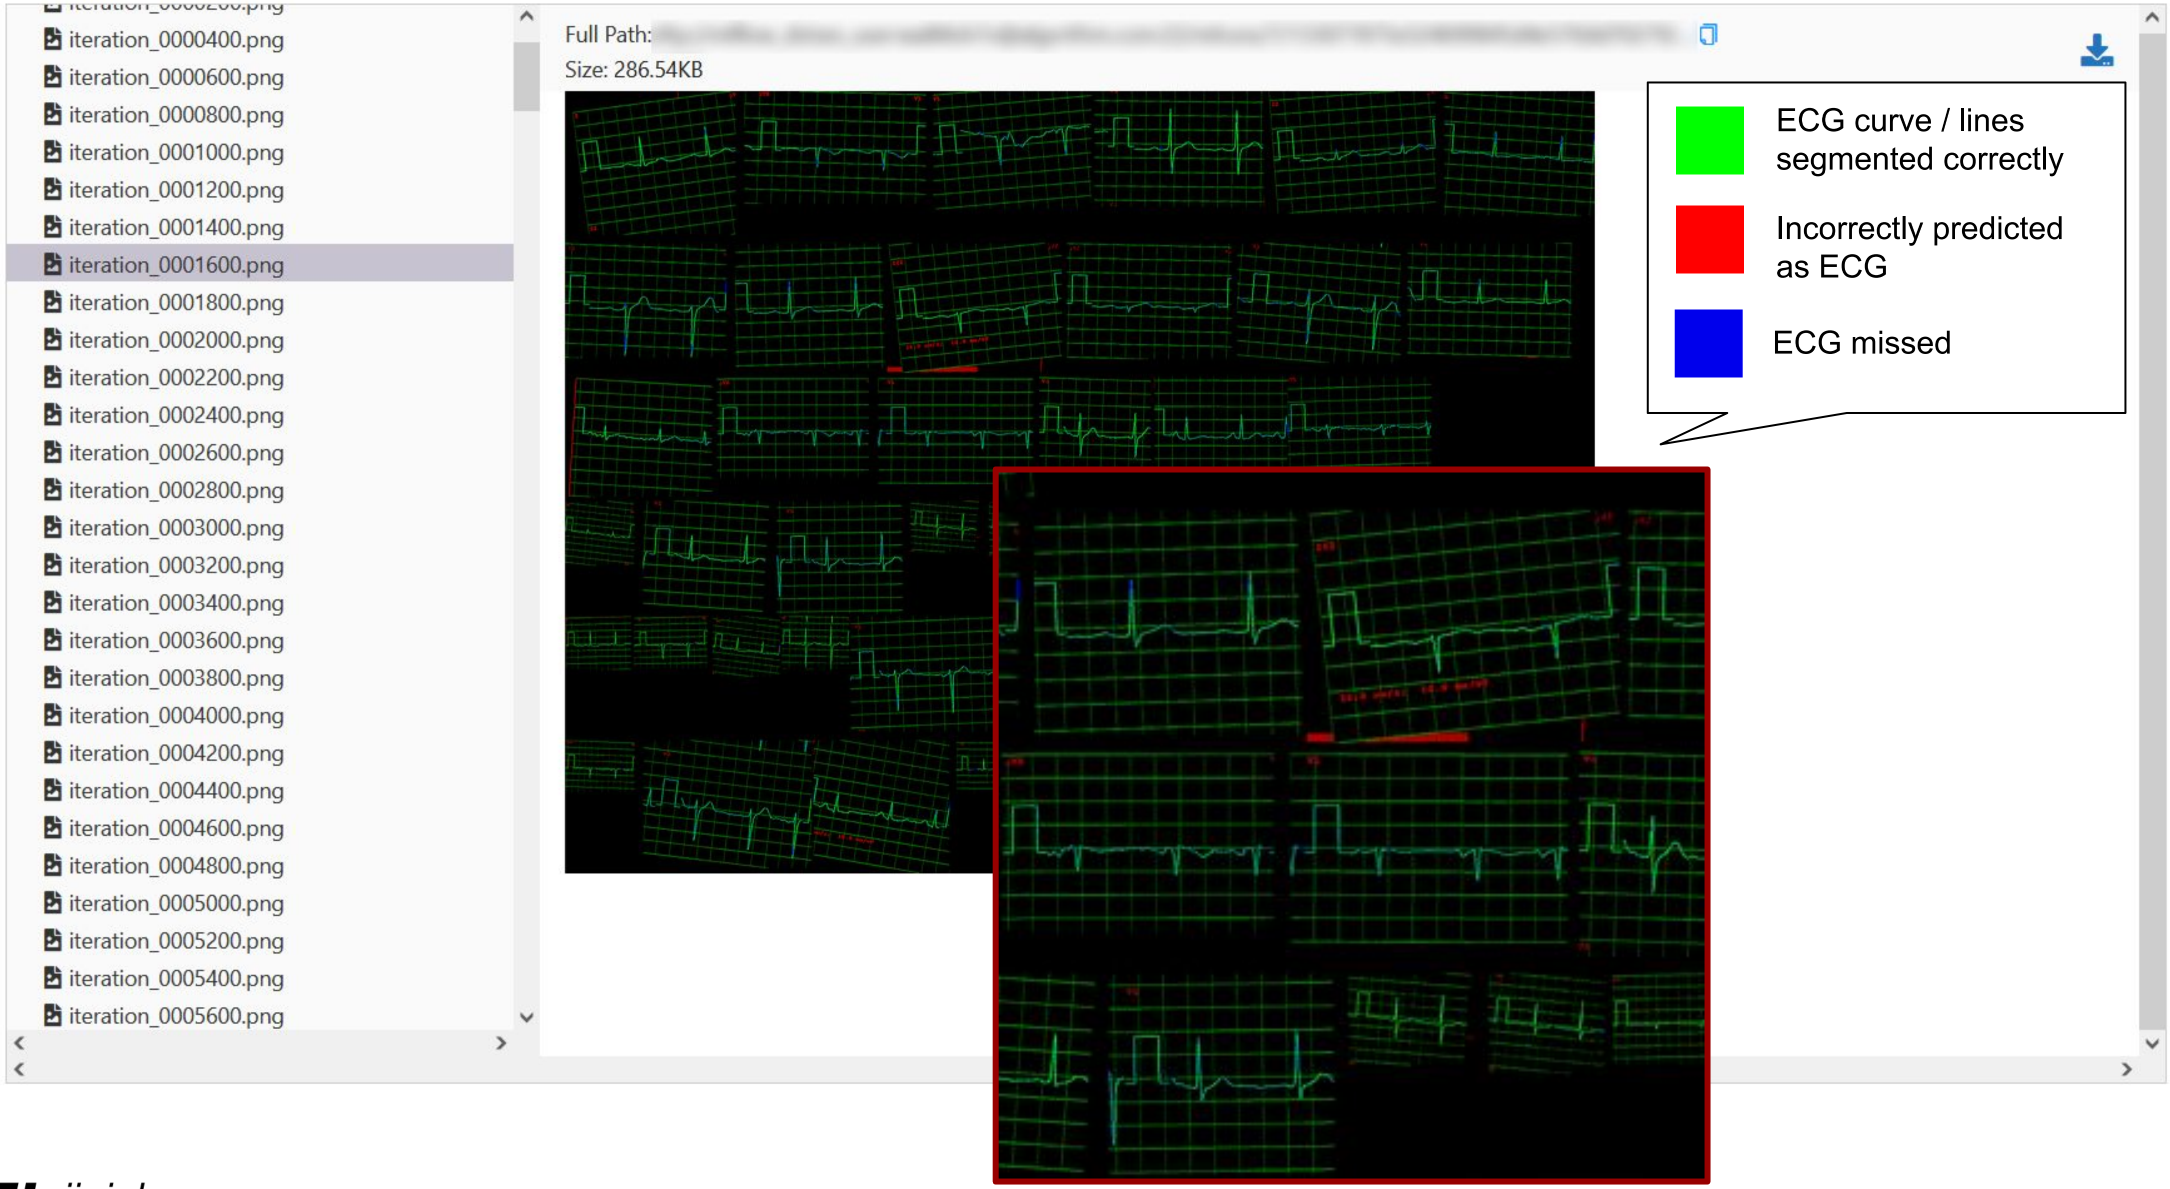

3.3 ECG Network Training: Training Frameworks

…after convergence, visualizations look as they should.

3.3 ECG Network Training: Metrics

Precision:

$$\frac{\#\ \text{True positives}}{\#\ \text{True positives} + \#\ \text{False positives}} \quad \in [0, 1]$$

i.e. "if the classifier predicts an MI, how likely is it really an MI?"

Recall:

$$\frac{\#\ \text{True positives}}{\#\ \text{True positives} + \#\ \text{False negatives}} \quad \in [0, 1]$$

i.e. "if there is an MI, how likely is the classifier to say so?"

F1 score:

$$\frac{2 \times \text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}} \quad \in [0, 1]$$

harmonic mean of precision and recall; means to combine the two scores

3.3 ECG Network Training: Results

Class	Precision	Recall	F1 Score
Normal ECG	74.59%	86.40%	0.8006
Non-MI-related abnormalities	69.50%	56.70%	0.6246
Myocardial Infarction	**61.20%**	**53.12%**	**0.5688**

Results on validation dataset of augmented ECG segmentation images

Validation dataset:
1331 ECG segmentations

581 normal ECGs
462 ECGs w/ **non-MI abnormalities**
288 ECGs w/ **MI**

3.3 ECG Network Training: Results

Class	Precision	Recall	F1 Score
Normal ECG	74.59%	86.40%	0.8006
Non-MI-related abnormalities	69.50%	56.70%	0.6246
Myocardial Infarction	**61.20%**	**53.12%**	**0.5688**

Results on validation dataset of augmented ECG segmentation images

Validation dataset:
1331 ECG segmentations

581 normal ECGs
462 ECGs w/ **non-MI abnormalities**
288 ECGs w/ **MI**

- ➢ Big problem: **<u>classification networks overfit!</u>**
  - ➢ Cannot perform data augmentation on <u>ECG segmentation images</u> easily on-the-fly
    - ➢ Need to **synthesize new ECG photographs**
    - ➢ **Time-consuming** process!

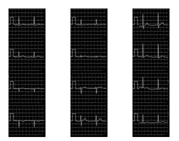

- ➢ One study (https://www.ncbi.nlm.nih.gov/pmc/articles/PMC8273385/) using **<u>10-second time series data</u>** (raw PTB-XL dataset) achieved **precision of 78.9%**, **recall of 81.8%**, **F1 score of 0.825 on MIs**
- ➢ We view our **result** as a **reasonable proof-of-concept**, considering we **only take snippets of ~3 seconds** for each lead and **use *distorted image data*** rather than pure time series, and did not perform hyperparameter tuning
  - ➢ **Shorter lead recordings** are much more **common in practice**

3.4 Web Tool

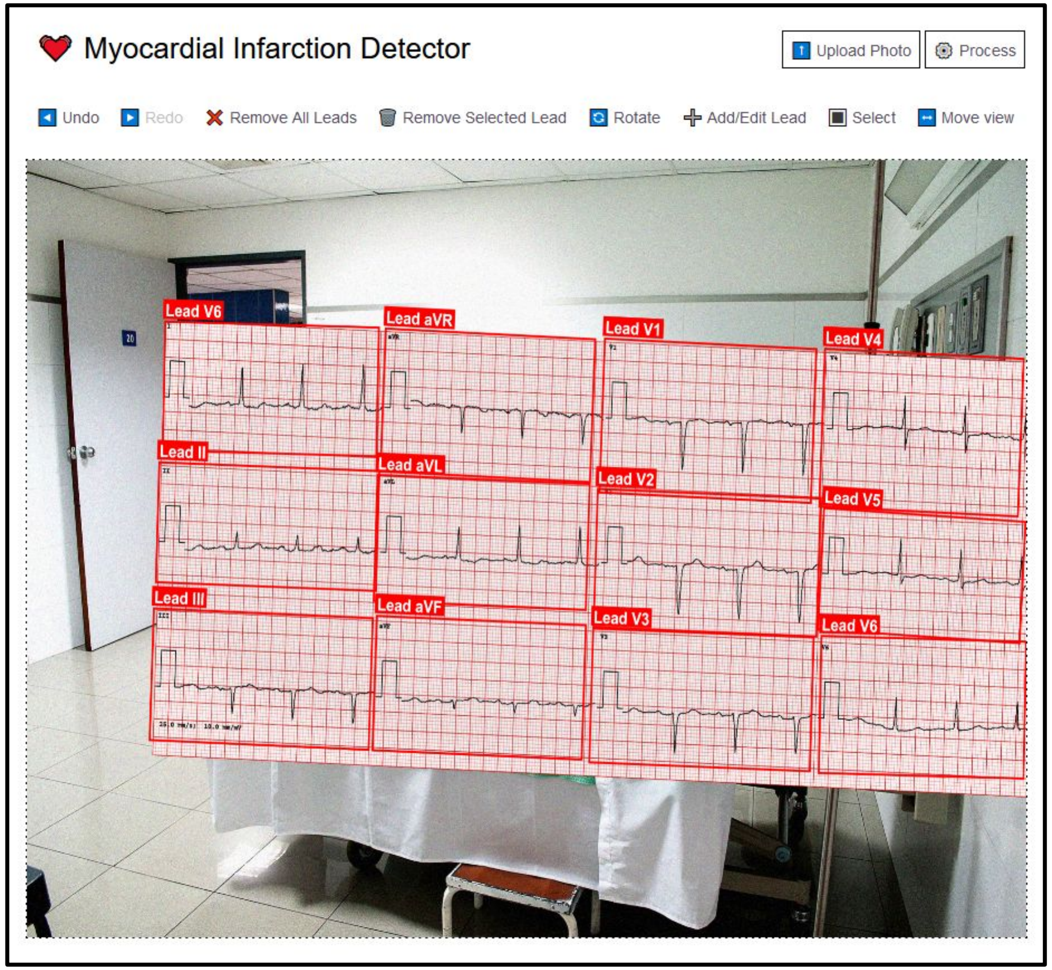

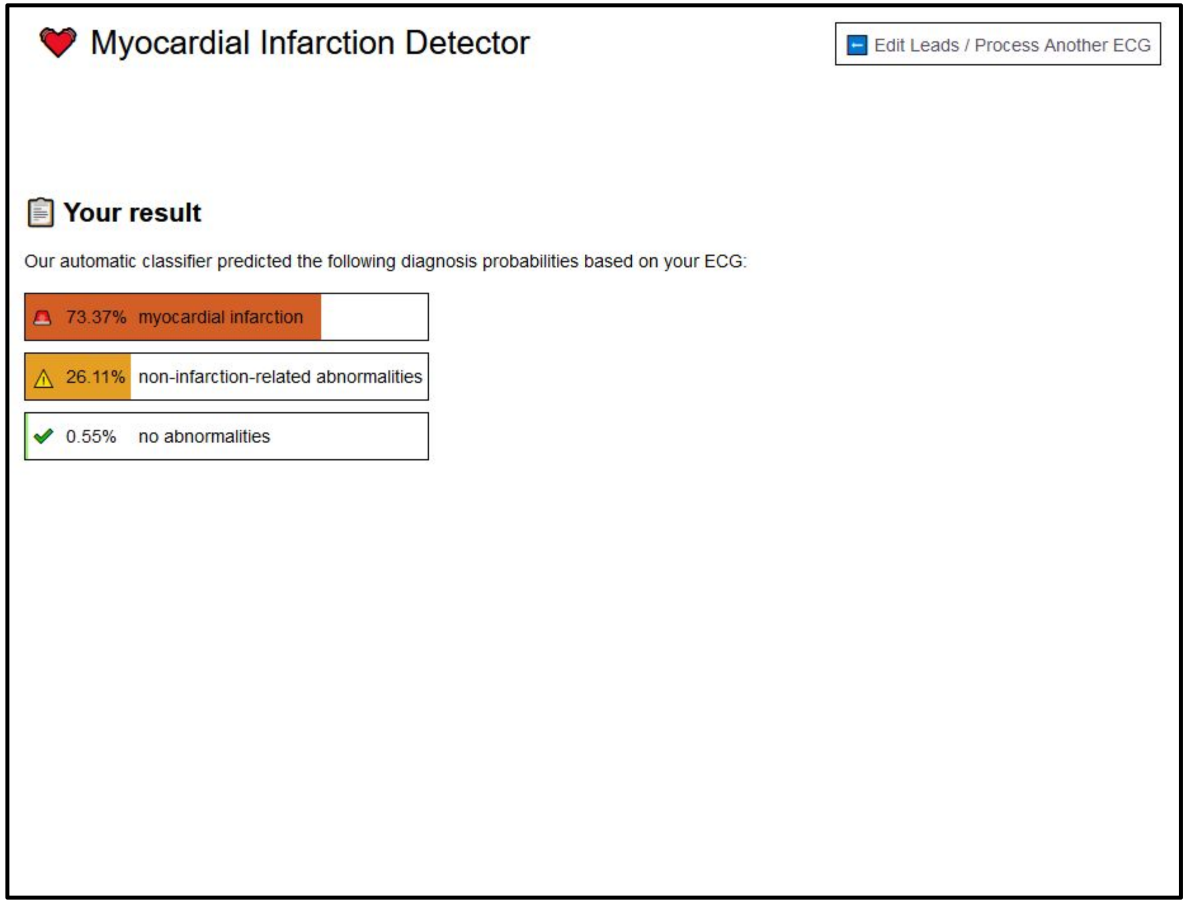

Current address (temporary): http://algvrithm.com:8080/

4. Future Work

- Use even more data!
  - Can generate nearly unlimited data, but: computationally expensive
  - Move image generation logic to GPU
- Diversify training set
  - Use different lead layouts & lengths
  - Use only subset of 12 leads
- Add support for paper speeds other than default 25mm/s
- Try other architectures for segmentation and classification
- Try to detect lead locations automatically, or provide initial suggestion for user to correct
- Implement noise removal algorithms
  - E.g. correction of baseline drift
- **Test using real ECGs!**

ETH zürich

DEMO

ETH zürich

THANK YOU!

💓

Q&A